


Parallel Programming with OpenMP



Copyright © 2009-2010 Yan Solihin

Copyright notice:

No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.



Module Outline

- Refresher on OpenMP
 - Model of parallelism
 - Components
 - Scope of variables
 - Synchronization
- OpenMP 3.0' s Tasking



Refresher on OpenMP

- Most popular shared memory programming standard
 - Backed by industry consortium
 - Open, not proprietary
 - Supported by most compilers, including GNU (starting from gcc version 4.2)
 - Still evolving (version 3.1 as of July 2011)
- Consists of directives, run-time system, and libraries

Refresher on OpenMP

- Initially, designed for expressing loop-level parallelism (i.e. parallelism between loop iterations)

Sequential Program

```
void main()
{
    int i, k, N=1000;
    double A[N], B[N], C[N];
    for (i=0; i<N; i++) {
        A[i] = B[i] + k*C[i]
    }
}
```

Parallel Program

```
#include "omp.h"
void main()
{
    int i, k, N=1000;
    double A[N], B[N], C[N];
    #pragma omp parallel for
    for (i=0; i<N; i++) {
        A[i] = B[i] + k*C[i];
    }
}
```

During Execution

- Single Program Multiple Data (SPMD)
 - The same code applied to different data

Thread 0

```
void main()
{
    int i, k, N;
    double A[N];
    lb = 0;
    ub = 250;
    for (i=lb; i<ub; i++) {
        A[i] = B[i] + k*C[i];
    }
}
```

Thread 1

```
#include "omp.h"
void main()
{
    int i, k, N=1000;
    double A[N], B[N], C[N];
    #pragma omp parallel for
    for (i=0; i<N; i++) {
        A[i] = B[i] + k*C[i];
    }
}
```

Thread 3

```
()
    k, N=1000;
    A[N], B[N], C[N];
    0;
    00;
    lb; i<ub; i++) {
        A[i] = B[i] + k*C[i];
    }
}
```

OpenMP Fork-and-Join model

```
printf("program begin\n");  
N = 1000;
```

Serial

```
#pragma omp parallel for  
for (i=0; i<N; i++)  
    A[i] = B[i] + C[i];
```

Parallel

```
M = 500;
```

Serial


```
#pragma omp parallel for  
for (j=0; j<M; j++)  
    p[j] = q[j] - r[j];
```

Parallel

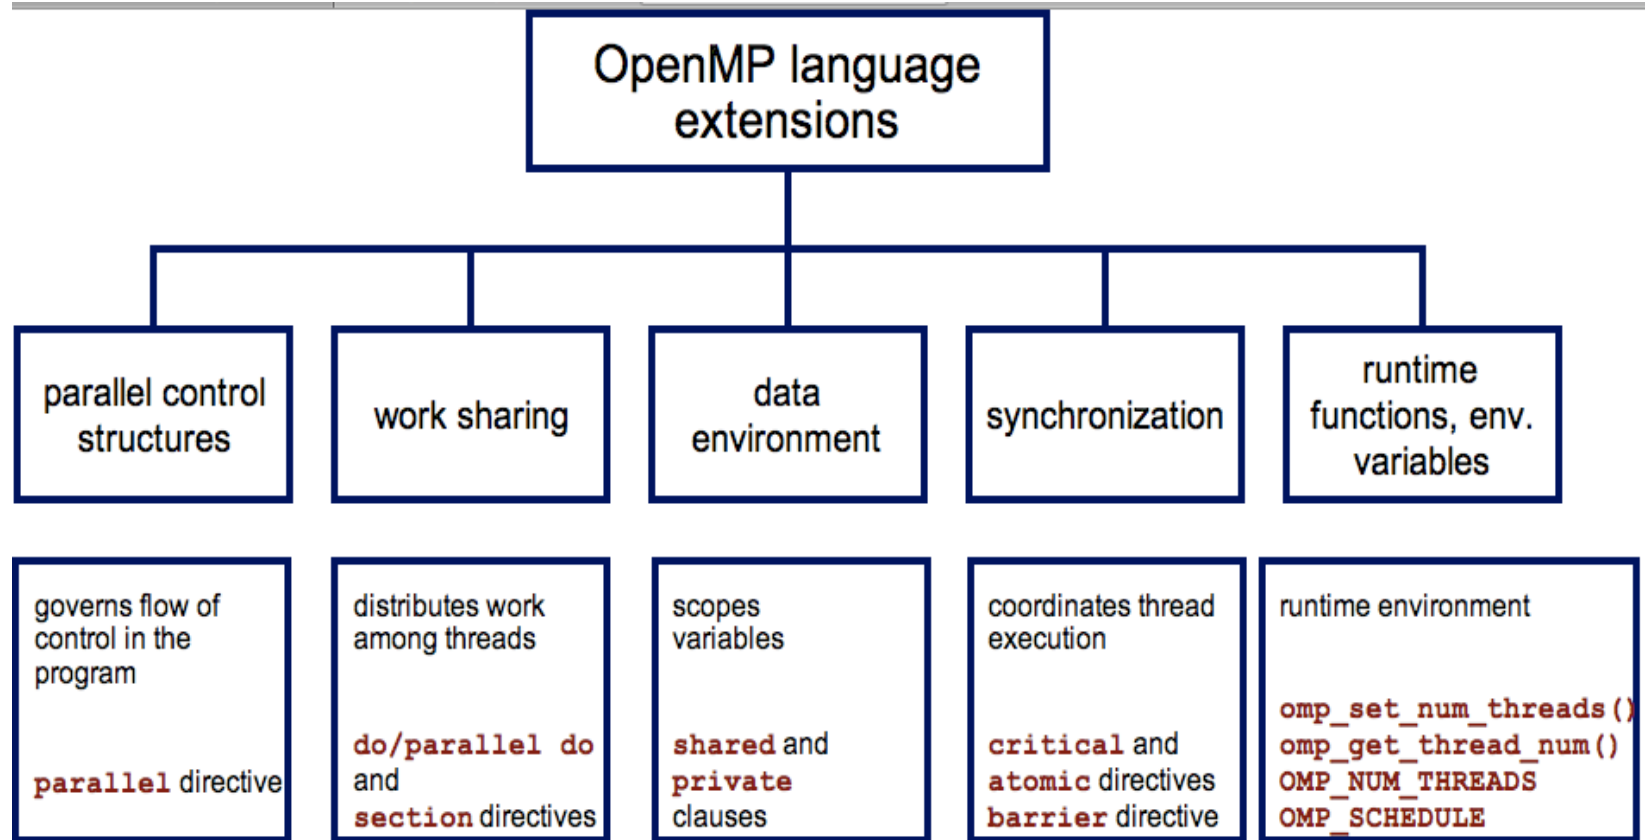
```
printf("program done\n");
```

Serial

Master thread



OpenMP's Components





Directives format

Refer to <http://www.openmp.org> and the OpenMP 2.0 specifications in the course web site for more details

```
#pragma omp directive-name [clause [ , clause ] ...] new-line
```

For example,

```
#pragma omp for [clause [ , clause ] ... ] new-line  
for-loop
```

The clause is one of

- **private**(*variable-list*)
- **firstprivate**(*variable-list*)
- **lastprivate**(*variable-list*)
- **reduction**(*operator*: *variable-list*)
- **ordered**
- **schedule**(*kind* [*,* *chunk_size*])
- **nowait**

Parallel for loop

■ Parallel for loop:

Threads spawned here ←

Code executed by
all threads

Threads joined here ←

```
#include <omp.h>
//...
#pragma omp parallel
{
    ...
    #pragma omp for private(i)
    for(i=0; i<n; i++)
        A[i] = A[i]*A[i] - 3.0;
}
```

Except here.
Iterations divided
over multiple threads

■ OpenMP allows us to express both the parallel region and work sharing of loop iterations by:

Threads spawned here ←
and iterations divided
over multiple threads

```
#pragma omp parallel for private(i)
for(i=0; i<n; i++)
    A[i] = A[i]*A[i] - 3.0;
```

Parallel Section

■ Parallel sections:

Threads spawned here ←

```
#pragma omp parallel shared(A,B) private(i)
```

Each section executed
by one thread ←

```
{  
    #pragma omp sections  
    {  
        #pragma omp section  
        for(i=0; i<n; i++)  
            A[i]= A[i]*A[i]- 4.0;  
        #pragma omp section  
        for(i=0; i<n; i++)  
            B[i]= B[i]*B[i] + 9.0;  
    } // end omp sections  
} // end omp parallel
```



Type of variables

- Most important: shared, private, reduction, firstprivate, lastprivate
- Semi-private data for parallel loops:
 - **reduction**: variable that is the target of a reduction operation performed by the loop, e.g., sum
 - **firstprivate**: initialize the private copy from the value of the shared variable prior to parallel section
 - **lastprivate**: upon loop exit, master thread holds the value seen by the thread assigned the last loop iteration (for parallel loops only)



Synchronization Constructs in OpenMP

- To enclose code in critical section, use `#pragma omp critical`

```
#include <omp.h>
//...
#pragma omp parallel
{
    ...
    #pragma omp parallel for private(i) {
        for(i=0; i<n; i++) {
            #pragma omp critical {
                sum = sum + A[i];
            }
        }
    }
}
```

- To enclose code that is to be executed by just one thread (the master thread), use `#pragma omp master`



Barriers

- Barriers are implicit after each parallel section
- When barriers are not needed for correctness, use **nowait** clause
- schedule clause will be discussed later

```
#include <omp.h>
//...
#pragma omp parallel
{
    ...
    #pragma omp parallel for nowait private(i)
        for(i=0; i<n; i++)
            A[i]= A[i]*A[i]- 3.0;
}
```



Other Synchronization Constructs

- `#pragma omp atomic` ensures critical section, but:
 - has lower overheads (implemented directly with machine instructions)
 - restricted to simple ops: adding/subtracting
 - e.g. `#pragma_omp_atomic`
 - `x = x+1;`
- Named locks
 - `omp_init_lock()`, `omp_set_lock()`, `omp_unset_lock()`, `omp_test_lock()`, `omp_destroy_lock()`
 - Use when you use multiple locks (fine-grain lock programming)



Declaring Variables

- Principle 1: declare variables as shared as much as possible
 - Read-only variables
 - A variable (or an element of a matrix) is only written and read by one thread
- Principle 2: variables that may be overwritten by another thread should be declared as private
 - Declaring a variable private creates replicas
 - e.g. `private(X)`
 - `X` (the original copy)
 - `X_private[0]` (replica for thread 0)
 - `X_private[1]` (replica for thread 1)
 - ...



Declaring Variables

- Principle 3: If a variable will aggregate values from all elements of an array or matrix, we can perform reduction
 - e.g. reduction(sum)

```
sum = 0;
#pragma omp parallel for reduction(+:sum)
for (i=0; i<N; i++) {
    sum = sum + a[i];
} // end omp parallel
```

- is equivalent to:

```
sum = mysum = 0;
#pragma omp parallel firstprivate(mysum) {
    #pragma omp for
    for (i=0; i<N; i++)
        mysum = mysum + a[i];
    #pragma omp critical {
        sum = sum + mysum
    }
} // end omp parallel
```




Restrictions on Reduction

- `reduction(op: var1, var2, ...)`
 - op can only be one of
 - arithmetic: `+`, `*`
 - logical: `&`, `|`, `&&`, `||`, `^`
 - mathematical: `min`, `max` (only in Fortran)



Declaring Variables

- Principle 4: some variables are treated as private by default
 - Variables declared inside the parallel region
 - Stack and local variables of a function
 - Function arguments are not private! (unless declared inside the parallel region)



Specifying Number of Threads

- Method 1: using environment variable, e.g.
 - setenv OMP_NUM_THREADS 2
 - a.out
 - In this case, OpenMP library spawns 2 threads to execute all parallel sections
- Method 2: hardwire it in the code

```
#include <omp.h>
...
omp_set_num_threads(4);
#pragma omp parallel // 4 threads will run
{
    // do work here
}

omp_set_num_threads(omp_num_procs());
#pragma omp parallel // as many threads as CPUs
{
    // do work here
}
```



Lab Assignment: Reduction

- Step 1: Go to Lab/Reduction
- Step 2: Parallelize the loop that reduces to sum
 - put `#pragma omp parallel for reduction(+:sum)`
- Step 3: Parallelize the loop that reduces to amax
 - declare “double my_amax;”
 - put `#pragma omp parallel for private(my_amax) {...}`
 - Replace “`amax = abs_max(amax,...)`” with “`my_amax = abs_max(my_amax, ...)`”. This causes each thread to accumulate its own amax value. We still need to reduce this at the end.
 - Add `#pragma omp critical {if (fabs(my_amax) > fabs(amax)) amax = my_amax;}` before exiting the parallel section
- Step 4: Compile with `gcc -fopenmp` flag.
- Step 5: Experiment with different number of threads
 - `setenv OMP_NUM_THREADS x`



OpenMP 3.0

- OpenMP version 3.0 added tasking ability (2008)
 - useful for irregular parallelism that is not loop-based and not easy to specify using sections
- **#pragma omp task** [*clause*[[,*clause*] ...] *structured-block*
- clause is one of
 - **if** (*expression*)
 - **untied**
 - **shared** (*list*)
 - **private** (*list*)
 - **firstprivate** (*list*)
 - **default(shared | none)**



Example of Tasking

- Suppose we are traversing a linked list and spawns a thread to process each node
 - we can let the master thread spawns a thread to process each node it visits

```
#pragma omp parallel {  
  #pragma omp single {  
    p = head;  
    while (p != NULL) {  
      #pragma omp task firstprivate(p) {  
        process (p);  
      }  
      p = p->next; ;  
    } // end while  
  }  
}
```

- must be careful to ensure task size > thread creation and management overheads