

Chapter 2:

Parallel Programming Models

Copyright @ 2005-2008 Yan Solihin

Copyright notice:

No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.



An exception is granted for academic lectures at universities and colleges, provided that the following text is included in such copy: "Source: Yan Solihin, Fundamentals of Parallel Computer Architecture, 2008".

Module 2.1 Parallel Programming Models

Programming Models

- What is programming model?
 - An abstraction provided by the hardware to programmers
 - Determines how easy/difficult for programmers to express their algorithms into computation tasks that the hardware understands
- Uniprocessor programming model
 - Based on program + data
 - Bundled in Instruction Set Architecture (ISA)
 - Highly successful in hiding hardware from programmers
- Multiprocessor programming model
 - Much debate, still searching for the right one...
 - Most popular: shared memory and message passing

Shared Mem vs. Msg Passing



- Shared Memory / Shared Address Space
 - Each memory location visible to all processors
- Message Passing
 - Each memory location visible to 1 processor

Thread/process – Uniproc analogy

Process: share nothing

```
if (fork() == 0)
    printf("I am the child process, my id is %d", getpid());
else
    printf("I am the parent process, my id is %d", getpid());
```

- heavyweight => high thread creation overhead
- The processes share nothing => explicit communication using socket, file, or messages

Thread: share everything

```
void sayhello() {
    printf("I am child thread, my id is %d", getpid());
}

printf("I am the parent thread, my id is %d", getpid());
clone(&sayhello, <stackarg>, <flags>, ())
```

- + lightweight => small thread creation overhead
- + The processes share addr space => implicit communication

Thread communication analogy

```
int a, b, signal;
...
void dosum(<args>) {
    while (signal == 0) {}; // wait until instructed to work
    printf("child thread> sum is %d", a + b);
    signal = 0; // my work is done
}

void main() {
    a = 5, b = 3;
    signal = 0;
    clone(&dosum,...) // spawn child thread
    signal = 1; // tell child to work
    while (signal == 1) {} // wait until child done
    printf("all done, exiting\n");
}
```

- Shared memory in multiproc provides similar memory sharing abstraction

Message Passing Example

```
Int a, b;
...
void dosum() {
    recvMsg(mainID, &a, &b);
    printf("child process> sum is %d", a + b);
}

Void main() {
    if (fork() == 0) // I am the child process
        dosum();
    else {           // I am the parent process
        a = 5, b = 3;
        sendMsg(childID, a, b);
        wait(childID);
        printf("all done, exiting\n");
    }
}
```

Differences with shared memory:

- Explicit communication
- Message send and receive provide automatic synchronization

Quantitative Comparison

Table 2.1: Comparing shared memory and message passing programming models.

Aspects	Shared Memory	Message Passing
Communication	implicit (via loads/stores)	explicit messages
Synchronization	explicit	implicit (via messages)
Hardware support	typically required	none
Development effort	lower	higher
Tuning effort	higher	lower

Development vs. Tuning Effort

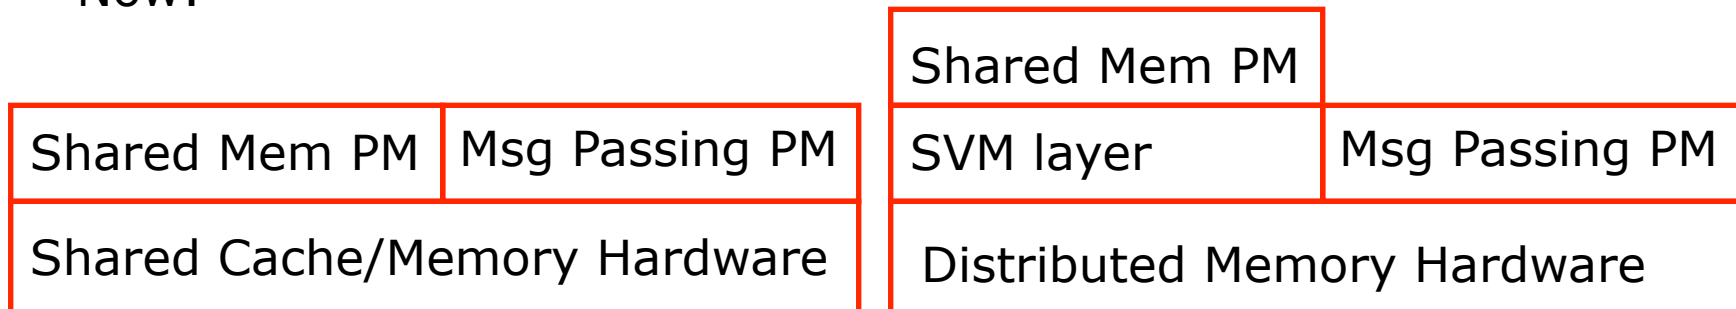
- Easier to develop shared memory programs
 - Transparent data layout
 - Transparent communication between processors
 - Code structure little changed
 - Parallelizing compiler, directive-driven compiler help
- Harder to tune shared memory programs for scalability
 - Data layout must be tuned
 - Communication pattern must be tuned
 - Machine topology matters for performance

Prog Model vs. Architecture

Was:



Now:



- Msg passing programs benefit from shared memory architecture
 - Sending a message achieved by passing a pointer to msg buffer
- Shared mem programs need software virtual memory (SVM) layer on distributed memory computers

More Shared Memory Example

```
for (i=0; i<8; i++)
    a[i] = b[i] + c[i];
sum = 0;
for (i=0; i<8; i++)
    if (a[i] > 0)
        sum = sum + a[i];
Print sum;
```

- + Communication directly through memory.
- + Requires less code modification
- Requires privatization prior to parallel execution

```
begin parallel // spawn a child thread
private int start_iter, end_iter, i;
shared int local_iter=4;
shared double sum=0.0, a[], b[], c[];
shared lock_type mylock;

start_iter = getid() * local_iter;
end_iter = start_iter + local_iter;
for (i=start_iter; i<end_iter; i++)
    a[i] = b[i] + c[i];
barrier;

for (i=start_iter; i<end_iter; i++)
    if (a[i] > 0) {
        lock(mylock) ;
        sum = sum + a[i];
        unlock(mylock) ;
    }
barrier;    // necessary

end parallel // kill the child thread
Print sum;
```

More Message Passing Example

```
for (i=0; i<8; i++)
    a[i] = b[i] + c[i];
sum = 0;
for (i=0; i<8; i++)
    if (a[i] > 0)
        sum = sum + a[i];
Print sum;
```

- + Communication only through messages
- Message sending and receiving overhead
- Requires algo and program modifications

```
// parent and child already spawned
id = getpid();
local_iter = 4;
start_iter = id * local_iter;
end_iter = start_iter + local_iter;

if (id == 0)
    send_msg (P1, b[4..7], c[4..7]);
else
    recv_msg (P0, &b[4..7], &c[4..7]);

for (i=start_iter; i<end_iter; i++)
    a[i] = b[i] + c[i];

local_sum = 0;
for (i=start_iter; i<end_iter; i++)
    if (a[i] > 0)
        local_sum = local_sum + a[i];
if (id == 0) {
    recv_msg (P1, &local_sum1);
    sum = local_sum + local_sum1;
    Print sum;
}
else
    send_msg (P0, local_sum);
```