

# Chapter 3

## Shared Memory Parallel Programming

Copyright @ 2005-2008 Yan Solihin

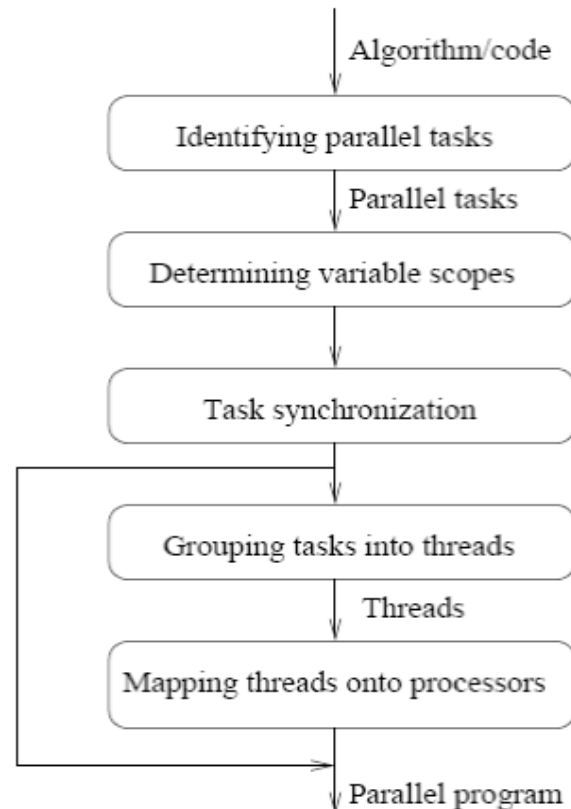
*Copyright notice:*

No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.

An exception is granted for academic lectures at universities and colleges, provided that the following text is included in such copy: "Source: Yan Solihin, Fundamentals of Parallel Computer Architecture, 2008".

# Module 3.1: Parallel Programming Techniques 1

# Steps in Creating a Parallel Program



Task Creation: identifying parallel tasks, variable scopes, synchronization  
Task Mapping: grouping tasks, mapping to processors/memory

# Parallel Programming

- Task Creation (correctness)
  - Finding parallel tasks
    - Code analysis
    - Algorithm analysis
  - Variable partitioning
    - Shared vs. private vs. reduction
  - Synchronization
- Task mapping (performance)
  - Static vs. dynamic
  - Block vs. cyclic
  - Dimension mapping: column-wise vs. row-wise
  - Communication and data locality considerations

# Code Analysis

- Goal: given a code, without the knowledge of the algorithms, find parallel tasks
- Focus on loop dependence analysis
- Notations:
  - $S$  is a statement in the source code
  - $S[i,j,\dots]$  denotes a statement in the loop iteration  $[i,j,\dots]$
  - $S1$  then  $S2$  means that  $S1$  *happens before*  $S2$
  - If  $S1$  then  $S2$ :
    - $S1 \rightarrow_T S2$  denotes true dependence, i.e.  $S1$  writes to a location that is read by  $S2$
    - $S1 \rightarrow_A S2$  denotes anti dependence, i.e.  $S1$  reads a location written by  $S2$
    - $S1 \rightarrow_O S2$  denotes output dependence, i.e.  $S1$  writes to the same location written by  $S2$

# Example

S1: $x = 2;$
S2: $y = x;$
S3: $y = x + z;$
S4: $z = 6;$

- Dependences:
  - $S1 \rightarrow_T S2$
  - $S1 \rightarrow_T S3$
  - $S3 \rightarrow_A S4$
  - $S2 \rightarrow_O S3$

# Loop-independent vs. loop-carried dependence

- Loop-carried dependence: dependence exists across iterations, i.e., if the loop is removed, the dependence no longer exists
- Loop-independent dependence: dependence exists within an iteration . i.e., if the loop is removed, the dependence exists

```
for (i=1; i<n; i++) {  
    S1: a[i] = a[i-1] + 1;  
    S2: b[i] = a[i];  
}  
  
for (i=1; i<n; i++)  
    for (j=1; j< n; j++)  
        S3: a[i][j] = a[i][j-1] + 1;  
  
for (i=1; i<n; i++)  
    for (j=1; j< n; j++)  
        S4: a[i][j] = a[i-1][j] + 1;
```

S1[i]  $\rightarrow$ T S1[i+1]: loop-carried  
S1[i]  $\rightarrow$ T S2[i]: loop-independent

S3[i,j]  $\rightarrow$ T S3[i,j+1]:  
-loop-carried on for j loop  
-No loop-carried dependence in  
for i loop

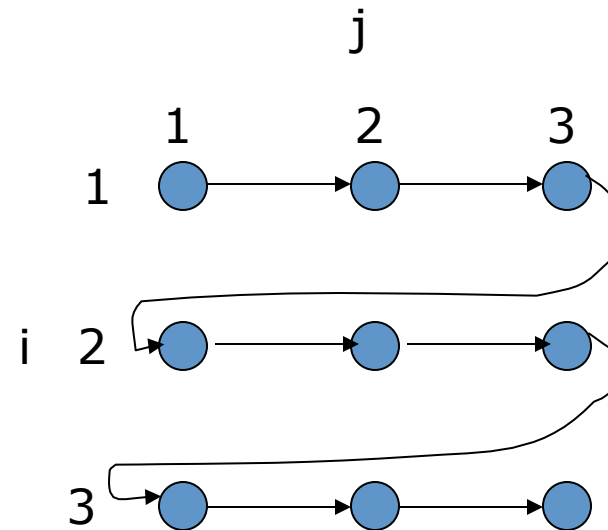
S4[i,j]  $\rightarrow$ T S4[i+1,j]:  
No loop-carried dependence in  
for j loop  
Loop-carried on for i loop

# Iteration-space Traversal Graph (ITG)

- ITG shows graphically the order of traversal in the iteration space (happens-before relationship)
- Node = a point in the iteration space
- Directed Edge = the next point that will be encountered after the current point is traversed

Example:

```
for (i=1; i<4; i++)  
  for (j=1; j<4; j++)  
    S3: a[i][j] = a[i][j-1] + 1;
```





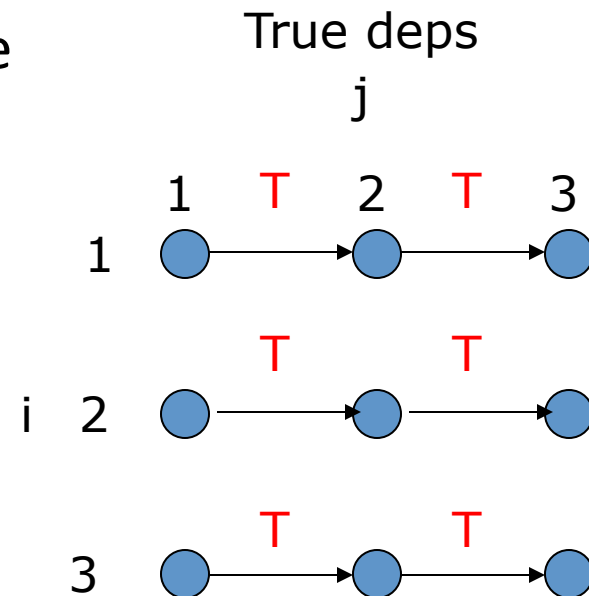
# Loop-carried Dependence Graph (LDG)

- LDG shows the true/anti/output dependence relationship graphically
- Node = a point in the iteration space
- Directed Edge = the dependence

Example:

```
for (i=1; i<4; i++)  
  for (j=1; j<4; j++)  
    S3: a[i][j] = a[i][j-1] + 1;
```

$S3[i,j] \rightarrow^T S3[i,j+1]$



# Further example

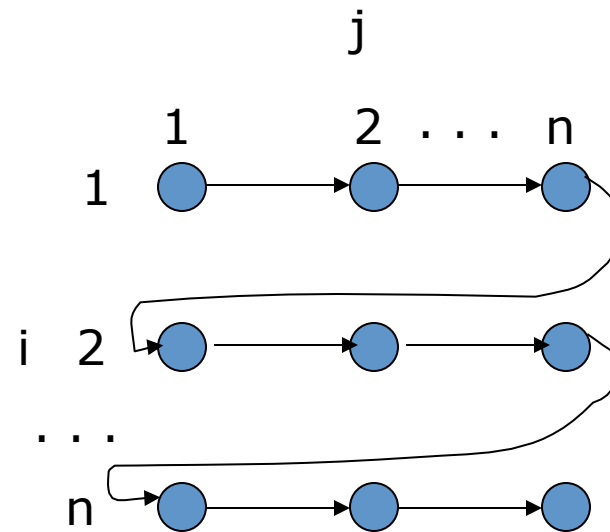
```
for (i=1; i<=n; i++)
  for (j=1; j<=n; j++)
    S1: a[i][j] = a[i][j-1] + a[i][j+1] + a[i-1][j] + a[i+1][j];

for (i=1; i<=n; i++)
  for (j=1; j<=n; j++) {
    S2: a[i][j] = b[i][j] + c[i][j];
    S3: b[i][j] = a[i][j-1] * d[i][j];
  }
```

- Draw the ITG
- List all the dependence relationships
- Draw the LDG

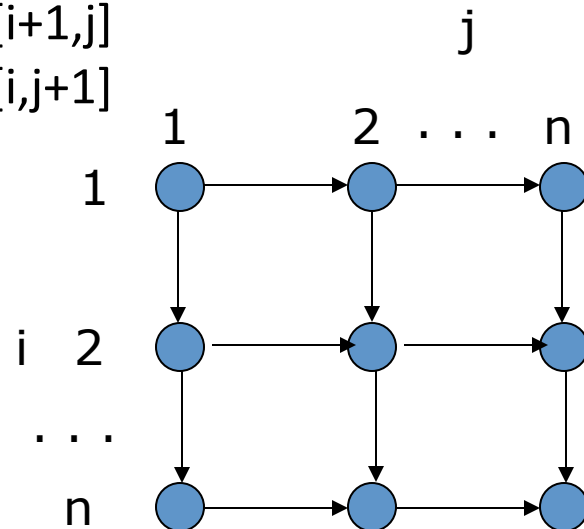
# Answer for Loop Nest 1

- ITG



# Answer for Loop Nest 1

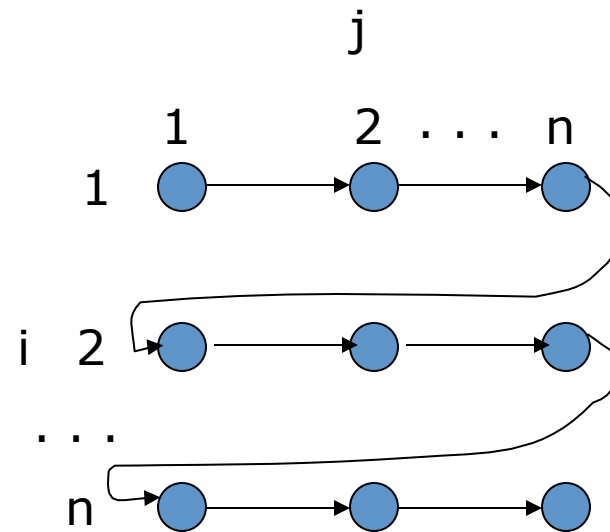
- True dependencies:
  - $S1[i,j] \rightarrow_T S1[i,j+1]$
  - $S1[i,j] \rightarrow_T S1[i+1,j]$
- Output dependencies:
  - None
- Anti dependencies:
  - $S1[i,j] \rightarrow_A S1[i+1,j]$
  - $S1[i,j] \rightarrow_A S1[i,j+1]$
- LDG:



Note: each Edge represents both true, and anti dependencies

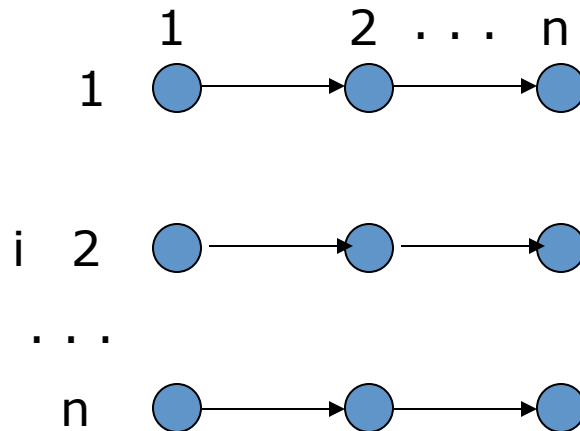
# Answer for Loop Nest 2

- ITG



# Answer for Loop Nest 2

- True dependences:
  - $S2[i,j] \rightarrow^T S3[i,j+1]$
- Output dependences:
  - None
- Anti dependences:
  - $S2[i,j] \rightarrow^A S3[i,j]$  (loop-independent dependence)
- LDG:



Note: each edge represents only true dependences

# Module 3.2: Parallel Programming Techniques 2

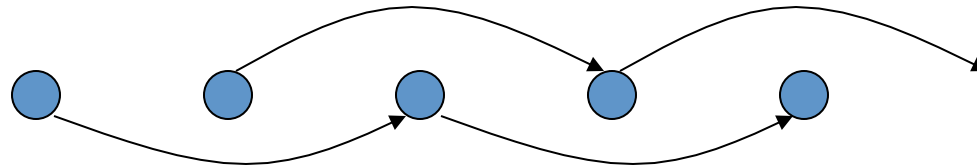
# Finding parallel tasks across iterations

- Analyze loop-carried dependencies:
  - Dependence must be obeyed (esp. true dependencies)
  - There are opportunities when some dependencies are missing

- Example 1:

```
for (i=2; i<=n; i++)  
  S: a[i] = a[i-2];
```

- LDG:



- Can divide the loop into two parallel tasks (one with odd iterations and another with even iterations):

```
for (i=2; i<=n; i+=2)  
  S: a[i] = a[i-2];  
for (i=3; i<=n; i+=2)  
  S: a[i] = a[i-2];
```

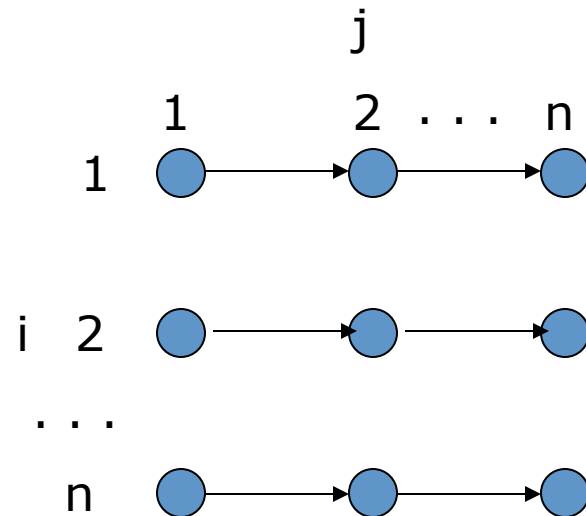


# Example 2

- Example 2:

```
for (i=0; i<n; i++)  
  for (j=0; j< n; j++)  
    S3: a[i][j] = a[i][j-1] + 1;
```

- LDG:

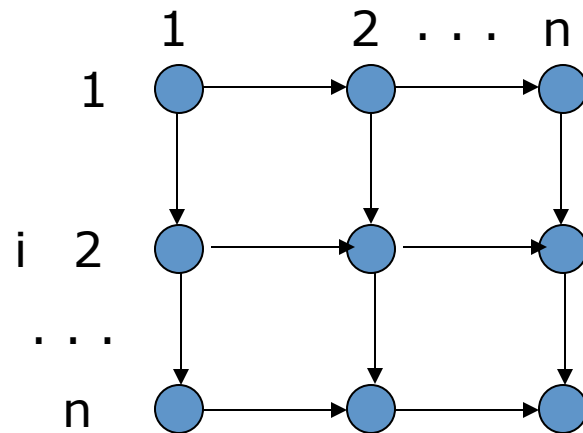


- There are n parallel tasks (one task per i iteration)

# Further example

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++)  
    S1: a[i][j] = a[i][j-1] + a[i][j+1] + a[i-1][j] + a[i+1][j];
```

- LDG:

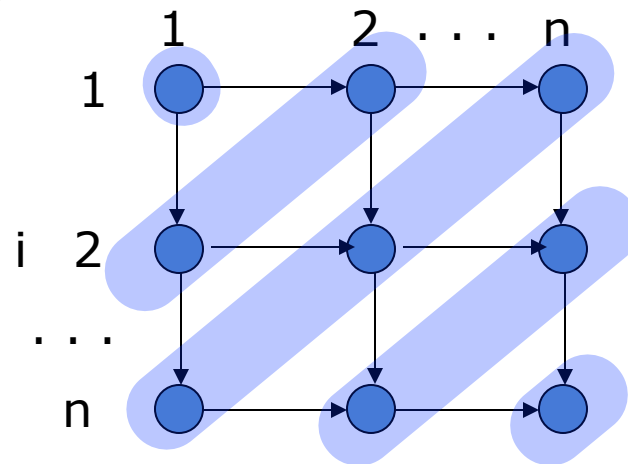


Note: each  
Edge represents  
both true, and  
anti dependences

- Where are the parallel tasks?

# Example 3

- Identify which nodes are not dependent on each other
- In each anti-diagonal, the nodes are independent of each other



Note: each  
Edge represents  
both true, and  
anti dependences

- Need to rewrite the code to iterate over anti-diagonals

# Structure of Rewritten Code

- Iterate over anti-diagonals, and over elements within an anti-diagonal:

```
Calculate number of anti-diagonals
Foreach anti-diagonal do:
    calculate number of points in the current anti-diagonal
    For each point in the current anti-diagonal do:
        compute the current point in the matrix
```

- Parallelize the highlighted loop
- Write the code...

# DOACROSS Parallelism

```
for (i=1; i<=N; i++) {  
    S: a[i] = a[i-1] + b[i] * c[i];  
}
```

Opportunity for parallelism?

$S[i] \rightarrow T S[i+1]$

So it has loop-carried dependence

Can change to:

```
for (i=1; i<=N; i++) {  
    S1: temp[i] = b[i] * c[i];  
}  
for (i=1; i<=N; i++) {  
    S2: a[i] = a[i-1] + temp[i];  
}
```

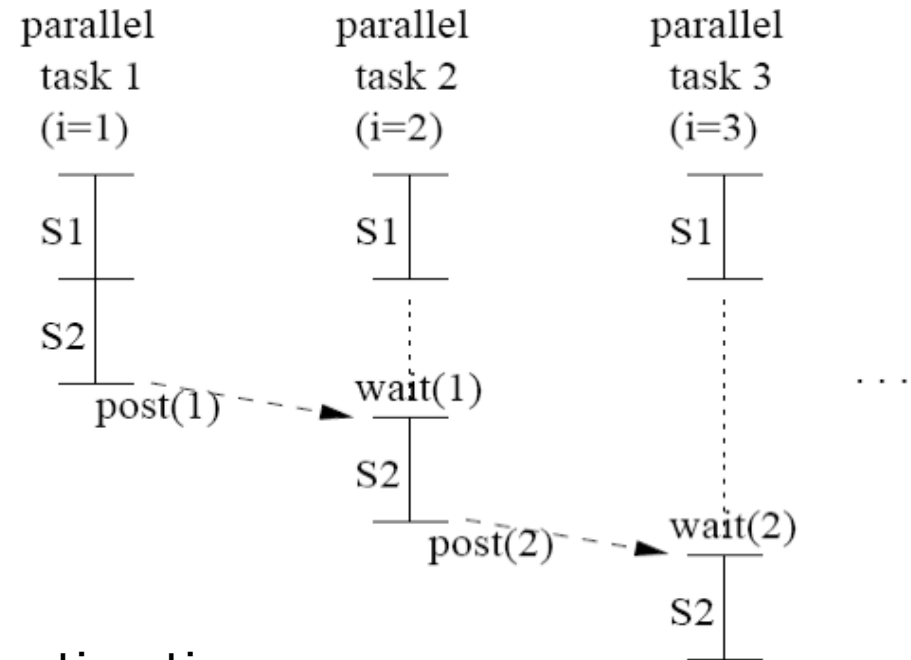
But, notice that the  $b[i] * c[i]$  part has no  
Loop-carried dependence

- Now the first loop is parallel, but the second one is not
- Execution time  $N \times (TS1 + TS2)$
- array temp[] introduces storage overhead
- Better solution?

# DOACROSS Parallelism

## DOACROSS Parallelism

```
Post(0);  
for (i=1; i<=N; i++) {  
    S1: temp = b[i] * c[i];  
    wait(i-1);  
    S2: a[i] = a[i-1] + temp;  
    post(i);  
}
```



Execution time now  
 $TS1 + N \times TS2$

Small storage overhead

# Finding Parallel Tasks in Loop Body

- Identify dependences in a loop body
- If there are independent statements, can split/distribute the loops

```
for (i=0; i<n; i++) {  
    S1: a[i] = b[i+1] * a[i-1];  
    S2: b[i] = b[i] * coef;  
    S3: c[i] = 0.5 * (c[i] + a[i]);  
    S4: d[i] = d[i-1] * d[i];  
}
```

Loop-carried dependences:

$S1[i] \rightarrow_A S2[i+1]$

Loop-indep dependences:

$S1[i] \rightarrow_T S3[i]$

- Note that S4 has no dependences with other statements
- “ $S1[i] \rightarrow_A S2[i+1]$ ” implies that S2 at iteration i+1 must be executed after S1 at iteration i. Hence dependence not violated if all S2's executed after all S1's

# After loop distribution

```
for (i=0; i<n; i++) {  
    S1: a[i] = b[i+1] * a[i-1];  
    S2: b[i] = b[i] * coef;  
    S3: c[i] = 0.5 * (c[i] + a[i]);  
    S4: d[i] = d[i-1] * d[i];  
}
```

- Each loop is a parallel task
- Referred to as **function parallelism**
- More distribution possible (refer to textbook)

```
for (i=0; i<n; i++) {  
    S1: a[i] = b[i+1] * a[i-1];  
    S2: b[i] = b[i] * coef;  
    S3: c[i] = 0.5 * (c[i] + a[i]);  
}  
  
for (i=0; i<n; i++) {  
    S4: d[i] = d[i-1] * d[i];  
}
```



# Identifying Concurrency (contd.)

- Function parallelism:
  - modest degree, does not grow with input size
  - ***difficult to load balance***
  - pipelining, as in video encoding/decoding, or polygon rendering
- Most scalable programs are data parallel
  - use both when data parallelism is limited

# DOPIPE Parallelism

```
for (i=2; i<=N; i++) {  
  S1: a[i] = a[i-1] + b[i];  
  S2: c[i] = c[i] + a[i];  
}
```

Loop-carried dependences:

$S1[i-1] \rightarrow T S1[i]$

Loop independent dependence:

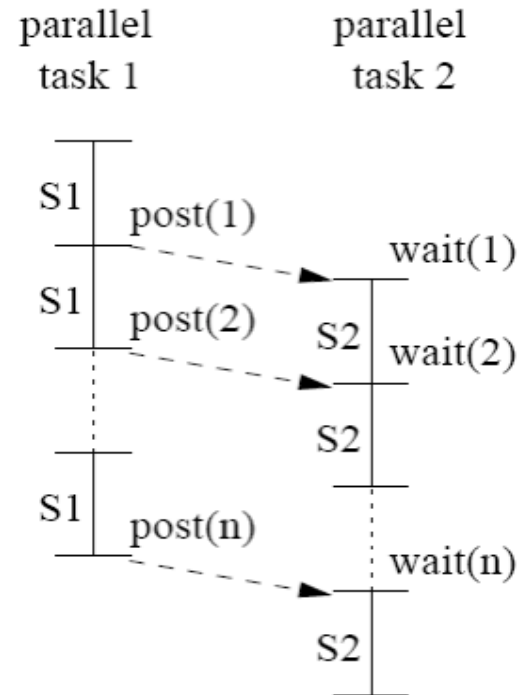
$S1[i] \rightarrow T S2[i]$

So, where is the parallelism opportunity?

## DOPIPE Parallelism

```
for (i=2; i<=N; i++) {  
  a[i] = a[i-1] + b[i];  
  post(i);  
}  
  
for (i=2; i<=N; i++) {  
  wait(i);  
  c[i] = c[i] + a[i];  
}
```

What is the max speedup?  
see textbook



# Parallel Programming

- Task Creation (correctness)
  - Finding parallel tasks
    - Code analysis
    - **Algorithm analysis**
  - Variable partitioning
    - Shared vs. private vs. reduction
  - Synchronization
- Task mapping (performance)
  - Static vs. dynamic
  - Block vs. cyclic
  - Dimension mapping: column-wise vs. row-wise
  - Communication and data locality considerations

# Task Creation: Algorithm Analysis

- Goal: code analysis misses parallelization opportunities available at the algorithm level
- Sometimes, the ITG introduces unnecessary serialization
- Consider the “ocean” algorithm
  - Numerical goal: at each sweep, compute how each point is affected by its neighbors
  - Hence, any order of update (within a sweep) is an approximation
  - Different ordering of updates: may converge quicker or slower
  - Change ordering to improve parallelism
  - Partition iteration space into red and black points
  - Red sweep and black sweep are each fully parallel

# Example 3: Simulating Ocean Currents

## Algorithm:

```
While not converging to a solution do:
  foreach timestep do:
    foreach cross section do a sweep:
      foreach point in a cross section do:
        compute the force interaction with its neighbors
```

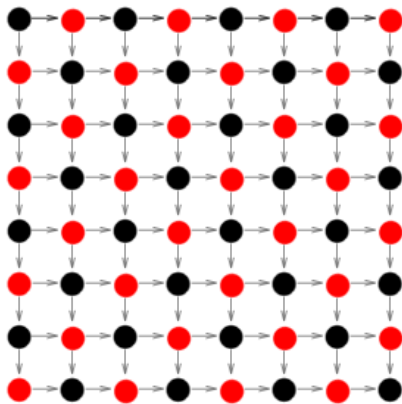
## compare with the code that implements the algorithm:

```
for (I=1; I<=N; I++) {
  for (j=1; j<=N; j++) {
    S1: temp = A[I][j];
    S2: A[I][j] = 0.2 * (A[I][j]+A[I][j-1]+A[I-1][j]+
                        +A[I][j+1]+A[I+1][j]);
    S3: diff += abs(A[I][j] - temp);
  }
}
```

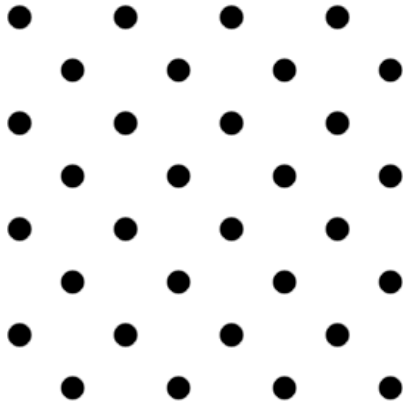
# Red-Black Coloring

in one sweep:  
- no dependence between  
black and red points

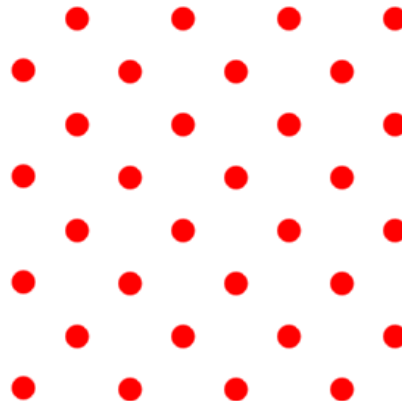
*LDG for a single sweep*



*LDG of black points (no dependences)*



*LDG of red points (no dependences)*



**restructured algorithm:**

```
While not converging to a solution do:  
  foreach timestep do:  
    foreach cross section do:  
      foreach red point do: //red sweep  
        compute the force interaction  
      wait until red sweep  
      foreach black point do: //blk sweep  
        compute the force interaction
```

**see textbook for code**

## Task Creation: Further Algorithm Analysis

- Can algorithm tolerate *asynchronous* execution?
  - simply ignore dependences within a sweep
  - parallel program *nondeterministic* (timing-dependent!)

```
for (I=1; I<=N; I++) {  
  for (j=1; j<=N; j++) {  
    S1: temp = A[I][j];  
    S2: A[I][j] = 0.2 * (A[I][j]+A[I][j-1]+A[I-1][j]+  
                        +A[I][j+1]+A[I+1][j]);  
    S3: diff += abs(A[I][j] - temp);  
  }  
}
```

# Module 3.3: Parallel Programming Techniques 3



# Parallel Programming

- Task Creation (correctness)
  - Finding parallel tasks
    - Code analysis
    - Algorithm analysis
  - Variable partitioning
    - Shared vs. private vs. reduction
  - Synchronization
- Task mapping (performance)
  - Static vs. dynamic
  - Block vs. cyclic
  - Dimension mapping: column-wise vs. row-wise
  - Communication and data locality considerations

# Determining Variable Scope

- This step is specific to shared memory programming model
- Analyze how each variable may be used across parallel tasks:
  - Read-only:
    - variable is only read by all tasks
  - R/W non-conflicting:
    - variable is read, written, or both by only one task
  - R/W Conflicting:
    - variable written by one task may be read by another

# Example 1

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S2: a[i][j] = b[i][j] + c[i][j];  
    S3: b[i][j] = a[i][j-1] * d[i][j];  
  }
```

- Define a parallel task as each “for i” loop iteration
- Read-only:
  - n, c, d
- R/W non-conflicting:
  - a, b
- R/W conflicting:
  - i, j

## Example 2

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S1: a[i][j] = b[i][j] + c[i][j];  
    S2: b[i][j] = a[i-1][j] * d[i][j];  
    S3: e[i][j] = a[i][j];  
  }
```

- Parallel task = each “for j” loop iteration
- Read-only:
  - n, i, c, d
- R/W Non-conflicting:
  - a, b, e
- R/W Conflicting:
  - j

# Privatization

- Privatization = converting a shared variable into a private variable in order to remove conflicts
  - Goal: R/W Conflicting → R/W Non-conflicting
- A conflicting variable is privatizable if
  - In program order, the variable is always defined (=written) by a task before use (=read) by the same task
  - The values for different parallel tasks are known ahead of time (hence, private copies can be initialized to the known values)
- Consequence
  - Conflicts disappear when the variable is “privatized”
- Privatization
  - involves making private copies of a shared variable
  - One private copy per *thread* (not per task)
  - How is this achieved in shared memory abstraction?

# Example 1

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S2: a[i][j] = b[i][j] + c[i][j];  
    S3: b[i][j] = a[i][j-1] * d[i][j];  
  }
```

- Define a parallel task as each “for i” loop iteration
- Read-only:
  - n, c, d
- R/W non-conflicting:
  - a, b
- R/W conflicting **but privatizable**:
  - i, j
  - After privatization: i[ID], j[ID]

## Example 2

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S1: a[i][j] = b[i][j] + c[i][j];  
    S2: b[i][j] = a[i-1][j] * d[i][j];  
    S3: e[i][j] = a[i][j];  
  }
```

- Parallel task = each “for j” loop iteration
- Read-only:
  - n, i, c, d
- R/W Non-conflicting:
  - a, b, e
- R/W Conflicting **but privatizable**:
  - j
  - After privatization: j[ID]

# Reduction

- Reduction
  - A special case of privatization, where:
  - Results are accumulated by each thread to a private copy
  - Private copies are merged into the shared copy at the end of computation
- Example: summing up array elements
  - Each thread works on its part of the array and accumulates its sum to a private sum
  - Private sums are accumulated into the shared sum at the end of computation



# Reduction Variables and Operations

- Reduction Operation examples:
  - SUM (+), multiplication (\*)
  - Logical (AND, OR, ...)
- Reduction variable =
  - The scalar variable that is the result of a reduction operation
- Criteria for reducibility:
  - Reduction variable is updated by each task, and the order of update is not important
  - Hence, the reduction operation must be **commutative** and **associative**

# Reduction Operation

- Compute:
  - $y = y\_init \ \underline{op} \ x1 \ \underline{op} \ x2 \ \underline{op} \ x3 \ \dots \ \underline{op} \ x_n$
- $op$  is a reduction operator if it is *commutative*
  - $u \ \underline{op} \ v = v \ \underline{op} \ u$
- and *associative*
  - $(u \ \underline{op} \ v) \ \underline{op} \ w = u \ \underline{op} \ (v \ \underline{op} \ w)$
- Certain operations can be transformed into reduction operations (see Homeworks)

# Variable Partitioning

- Should be declared private:
  - Privatizable variables
- Should be declared shared:
  - Read-only variables
  - R/W Non-conflicting variables
- Should be declared reduction:
  - Reduction variables
- Other R/W Conflicting variables:
  - Privatization possible? If so, privatize them
  - Otherwise, declare as shared, but protect with synchronization

# Example 1

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S2: a[i][j] = b[i][j] + c[i][j];  
    S3: b[i][j] = a[i][j-1] * d[i][j];  
  }
```

- Declare as shared:
  - n, c, d, a, b
- Declare as private:
  - i, j

## Example 2

```
for (i=1; i<=n; i++)  
  for (j=1; j<=n; j++) {  
    S1: a[i][j] = b[i][j] + c[i][j];  
    S2: b[i][j] = a[i-1][j] * d[i][j];  
    S3: e[i][j] = a[i][j];  
  }
```

- Parallel task = each “for j” loop iteration
- Declare as shared:
  - n, i, c, d, a, b, e
- Declare as private:
  - j

# Module 3.4: Parallel Programming Techniques 4

# Parallel Programming

- Task Creation (correctness)
  - Finding parallel tasks
    - Code analysis
    - Algorithm analysis
  - Variable partitioning
    - Shared vs. private vs. reduction
  - **Synchronization**
- Task mapping (performance)
  - Static vs. dynamic
  - Block vs. cyclic
  - Dimension mapping: column-wise vs. row-wise
  - Communication and data locality considerations

# Synchronization Primitives

- Point-to-point
  - a pair of `post()` and `wait()`
  - a pair of `send()` and `recv()` in message passing
- Lock
  - ensures mutual exclusion, only one thread can be in a locked region at a given time
- Barrier
  - a point where a thread is allowed to go past it only when all threads have reached the point.



# Lock

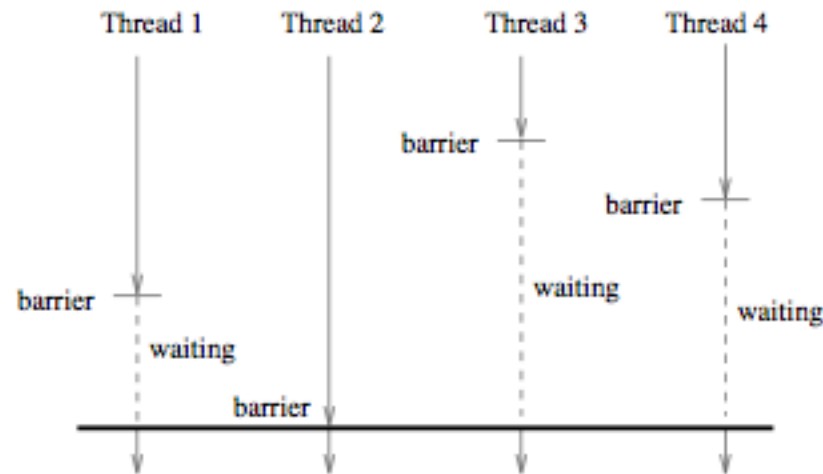
- What problem may arise here?

```
// inside a parallel region
for (i=start_iter; i<end_iter; i++)
    sum = sum + a[i];
```

- Lock ensures only one thread inside the locked region

```
// inside a parallel region
for (i=start_iter; i<end_iter; i++) {
    lock(x);
    sum = sum + a[i];
    unlock(x);
}
```

# Barrier: Global Event Synchronization



- Load balance important
- Execution time dependent on the slowest thread
  - One reason for gang scheduling and avoiding time sharing and context switching