



Chapter 5

Parallel Programming for Linked Data Structures

Copyright © 2005-2008 Yan Solihin

Copyright notice:

No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.

An exception is granted for academic lectures at universities and colleges, provided that the following text is included in such copy: "Source: Yan Solihin, Fundamentals of Parallel Computer Architecture, 2008".

Linked Data Structures

- Examples: linked lists, trees, graphs, hash tables
- Commonly used in many non-numerical programs
- Common features:
 - Operations involve: node insertion, node deletion, node search
 - Traversal follows pointer chasing pattern -> loop-carried dependence
- Example:

```
void addValue(pIntList pList, int key, int x)
{
    pIntListNode p = pList->head;
    while (p != NULL) {
        if (p->key == key)
            p->data = p->data + x;
        p = p->next;
    }
}
```

How to Parallelize LDS

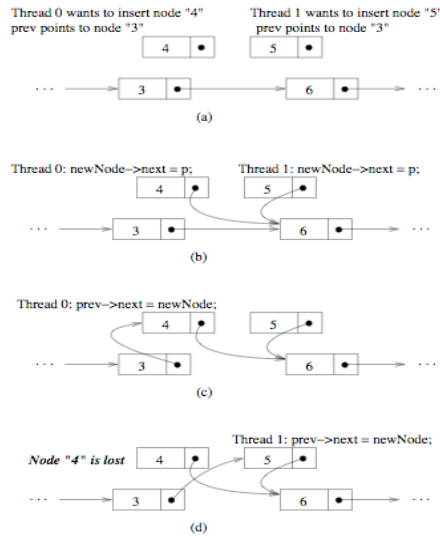
- Loop-level analysis does not reveal parallelism
- Look at algorithm level
- LDS is a data structure with operations that can be performed on them
 - node insertion
 - node deletion
 - node search
 - etc.
- Conceptually, we can allow several operations to be performed in parallel
- But how do we reason about the correctness of parallel operations?

Correctness of Parallel LDS Operations

- Serializability = *a parallel execution of a group of operations or primitives is serializable if there is some serial primitives that produce an identical result*
- Suppose a node insertion and node deletion are performed in parallel
- The outcome must be equivalent to either
 - node insertion is performed after node deletion, or
 - node deletion is performed after node insertion
 - nothing else can be considered correct
- We will take a look at a simple case of a singly-linked list

Serializability between 2 Insertions

Conflict between 2 insertion operations



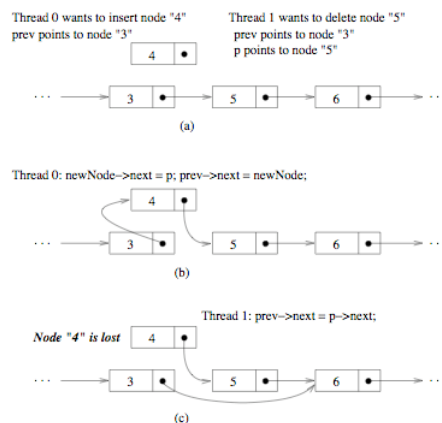
- Serializable outcome:
 - insert 4, then insert 5
 - or insert 5, then insert 4
 - in both cases, both nodes 4 and 5 must be in the list at the end of execution
- Incorrect parallel execution may result in:
 - node 4 lost

Fundamentals of Parallel Computer Architecture - Chapter 5

5

Serializability b/w Insertion and Deletion

Conflict between an insertion and a deletion operations



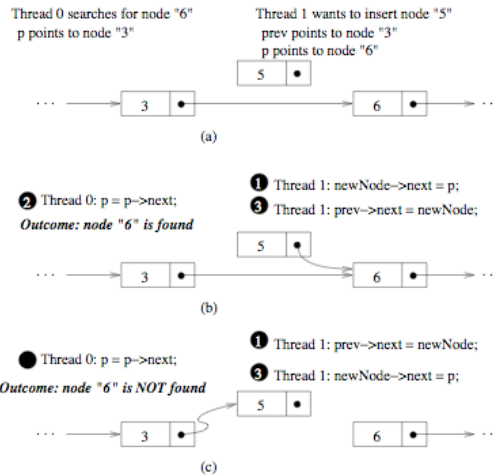
- Serializable outcome
 - insert 4, then delete 5
 - or delete 5, then insert 4
 - in both cases, at the end of execution, node 4 is in the list, node 5 is not in the list
- Incorrect parallel execution:
 - node 4 is lost, never inserted

Fundamentals of Parallel Computer Architecture - Chapter 5

6

Serializability b/w Insertion and Search

Conflict between an insertion and a search operations



Insertion operation and node search operations


- insert 5, then search 6
- or search 6, then insert 5
- in both cases, at the end of execution, 5 must be in the list, and 6 is found

Incorrect parallel execution

- node 6 may be found
- but node 6 may also not be found, and the linked list is broken

Serializability between Deletion and Search

Conflict between a deletion and a search operations



Deletion and search

- delete 5, then search 5
- search 5, then delete 5

Outcome is serializable

- Node 5 may be found or not found
- Node 5 is deleted from the list

Main Observations

- *Parallel execution of a two operations that affect a common node, in which at least one operation involves writing to the node, can produce conflicts that lead to non-serializable outcome.*
- *Under some circumstances, serializable outcome may still be achieved under a conflict mentioned in the above point.*
- *Conflicts can also occur between LDS operations and memory management functions such as memory deallocation and allocation.*

Parallelization Strategies

- Parallelization among readers
 - Very simple
 - But only parallel if there are many readers
- Global lock approach
 - Relatively simple
 - Parallel traversal, followed by sequential list modifications
- Fine-grain lock approach
 - Each node is associated with a lock
 - Each operation locks only nodes that need to be exclusively accessed
 - Complex: deadlock can occur, memory allocation and deallocation more complex

Parallelization Among Readers

- Basic idea:
 - (read only) operations that do not modify the list can go in parallel
 - (write) operations that modify the list execute sequentially
- How to enforce:
 - a read-only operation acquires a read lock
 - a write operation acquires a write lock
- Construct a lock compatibility table

Already Granted Lock	Read Lock requested	Write Lock requested
Read Lock	Yes	No
Write Lock	No	No

Fundamentals of Parallel Computer Architecture - Chapter 5

11

Example

```
IntListNode_Search(int x)
{
    acq_read_lock();
    ...
    ...
    rel_read_lock();
}
```

```
IntListNode_Insert(node *p)
{
    acq_write_lock();
    ...
    ...
    rel_write_lock();
}
```

Fundamentals of Parallel Computer Architecture - Chapter 5

12

Global Lock Approach

- Each operation logically has two steps
 - Traversal
 - Node insertion: find the correct location for the node
 - Node deletion: find the node to delete
 - Node search: find the node in question
 - List modification
- Basic idea: perform the traversal in parallel, but list modification in a critical section
- Pitfall:
 - prior and after entering critical section, the list may have changed
 - so the assumptions must be validated

Example

```
IntListNode_Search(int x)          IntListNode_Insert(node *p)
{
    acq_read_lock();
    ...
    ...
    ...
    rel_read_lock();
}

/* perform traversal */
/* with read locking */
...
acq_write_lock();
...
/* modifies list */
...
rel_write_lock();
}
```

Fine-Grain Locking Approach

- Associate each node with a lock (read, write)
- Each operation locks only needed nodes
- (Read and write) operations execute in parallel except when they conflict on some nodes
- Nodes that will be modified are write locked
- Nodes that are read and must remain unchanged are read locked
- Pitfall: deadlocks becomes possible
- Deadlocks can be avoided by imposing a global lock acquisition order

Example

- Refer to code example in the book