

Chapter 11

Distributed Shared Memory

Multiprocessors

Copyright © 2005-2008 Yan Solihin

Copyright notice:

No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.

An exception is granted for academic lectures at universities and colleges, provided that the following text is included in such copy: "Source: Yan Solihin, Fundamentals of Parallel Computer Architecture, 2008".

Module 11.1 – Basic Directory Protocol

Scalable Shared Memory System

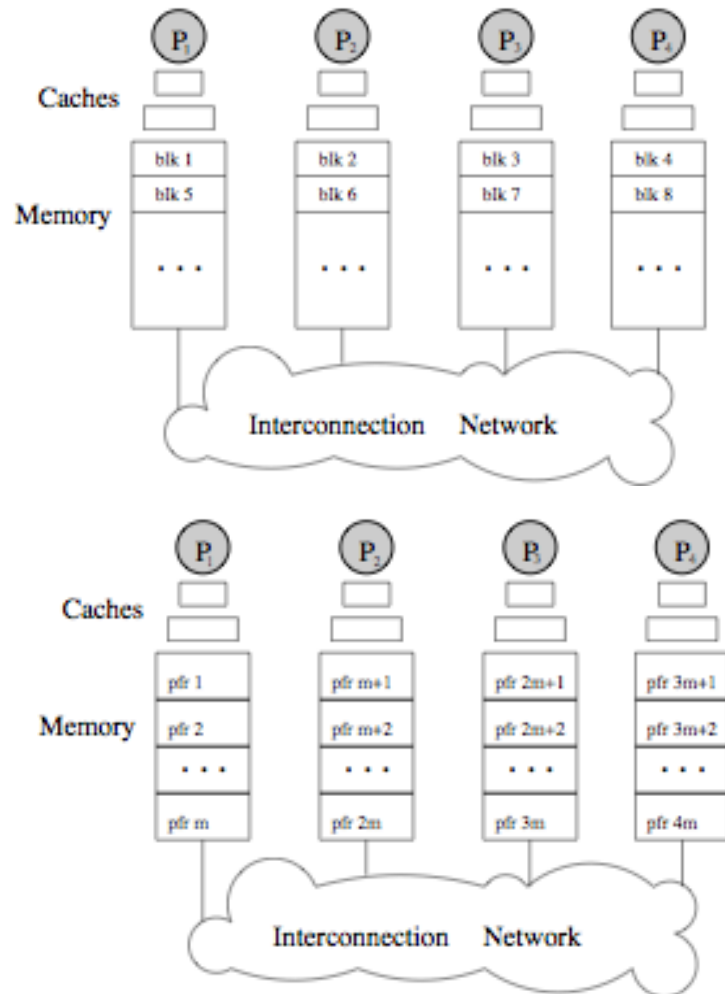
- Bus-based MP: Good for small multiprocessors, but does not scale
 - Physical constraints
 - long wires (low clock frequency), arbitration delay
 - Protocol constraints
 - Snoopy/broadcast: bandwidth getting saturated quickly
 - Contention everywhere: bus, snoopers, memory
- How to scale to larger MP?

Ways to Scale Bus-Based MP

		<i>Interconnection</i>	
		<i>Bus</i>	<i>Point-to-point</i>
<i>Protocol</i>	<i>Snoopy</i>	Least scalable	More scalable
	<i>Directory</i>		Most scalable

- Approach 1: replace bus with point-to-point interconnection, but keep relying on snooping/broadcast
 - e.g., AMD Opteron (mesh), IBM Cell (ring), Intel Larrabee (ring)
 - Still cannot scale to hundreds of processors
- Approach 2: replace broadcast with directory protocol
 - e.g. SGI Altix system
 - Leads to distributed shared memory (DSM) multiprocessors
 - Scaling to tens to hundreds of processors

How to Map Memory on DSM?




- Block interleaving?
 - distributes data around
 - Hard to exploit locality
- No interleaving?
 - OS must be involved in deciding where to allocate a page
 - with significant impact on performance
 - First touch vs. round robin

Finding a Block


- A block can only map to a single memory
- This memory is referred to as the “home” of the block
- The home memory keeps track of which caches may have the copy of a block
 - in a structured called the “directory”
- Where is directory stored?
 - Memory (entirely) vs. Memory+cache
 - Centralized vs. distributed

Basic Directory Handling of a Load Request



- 1: Requestor (P_1) sends a Read request to the Home (P_2) of the block
- 2: The home looks up the directory to find out who has the block and in what state. Then it sends an intervention request to the current owner (P_3)
- 3: The owner (P_3) supplies the block to the requestor (P_1) and may update the home as well

Basic Directory Handling of a Write Req



- 1: Requestor (P_1) sends a Write request to Home (P_2)
- 2: Home finds out that P_2 and P_3 are currently sharers, then it sends invalidation messages to them
- 3: Sharers (P_2 and P_3) reply with Invalidation Acknowledgement to P_1

What Info in the Directory

- Information that must be kept at the directory
 - Which caches are currently sharers or owner
 - Which state the block is currently cached
- Ideally, if the caches support MESI states, then the directory keeps MESI states as well
- But:
 - Transition from E to M in a cache is silent
 - Hence, the directory cannot tell if a block is cached in E or M state
- Hence, the directory may keep 3 states
 - E/M: if the block is cached in E or M state
 - S: if the block is cached in S state
 - U: if the block is uncached or cached in invalid state

Performance Criteria

- Main criteria:
 - **Traffic on invalidations**
 - **Latency of invalidations**
 - **Storage overheads**
- Less important criteria:
 - Traffic on interventions
 - Latency of interventions
- Invalidation involves multiple sharers vs. intervention involves a single owner

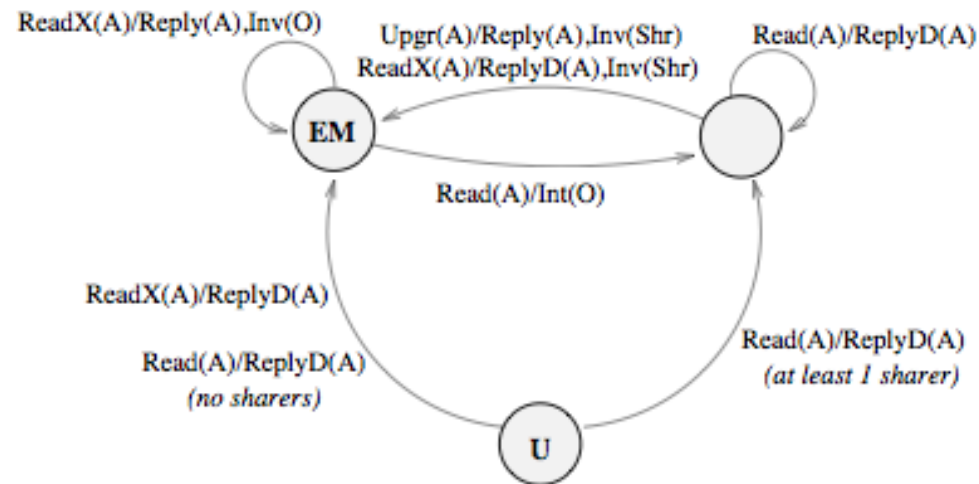
Example

- Cache State: MESI
- Directory:
 - Full-bit Vector
 - 3 States:
 - EM (Exclusive or Modified)
 - S (Shared)
 - U (Unowned)

Full Bit-Vector: Example

- Coherence messages
 - Issued by requesting processor:
 - Read: read request
 - ReadX: read exclusive (or write) request
 - Upgr: upgrade request
 - Issued by home:
 - ReplyD: home replies with data to requestor
 - Reply: home replies not containing data
 - Inv: home asking sharer to invalidate
 - Int: home asking owner to flush and change to S
 - Issued by owner of a cache block:
 - Flush: owner flushes data to home + requestor
 - InvAck: sharer/owner ack Invalidation msg
 - Ack: acknowledgement of receipt of non-inval msg

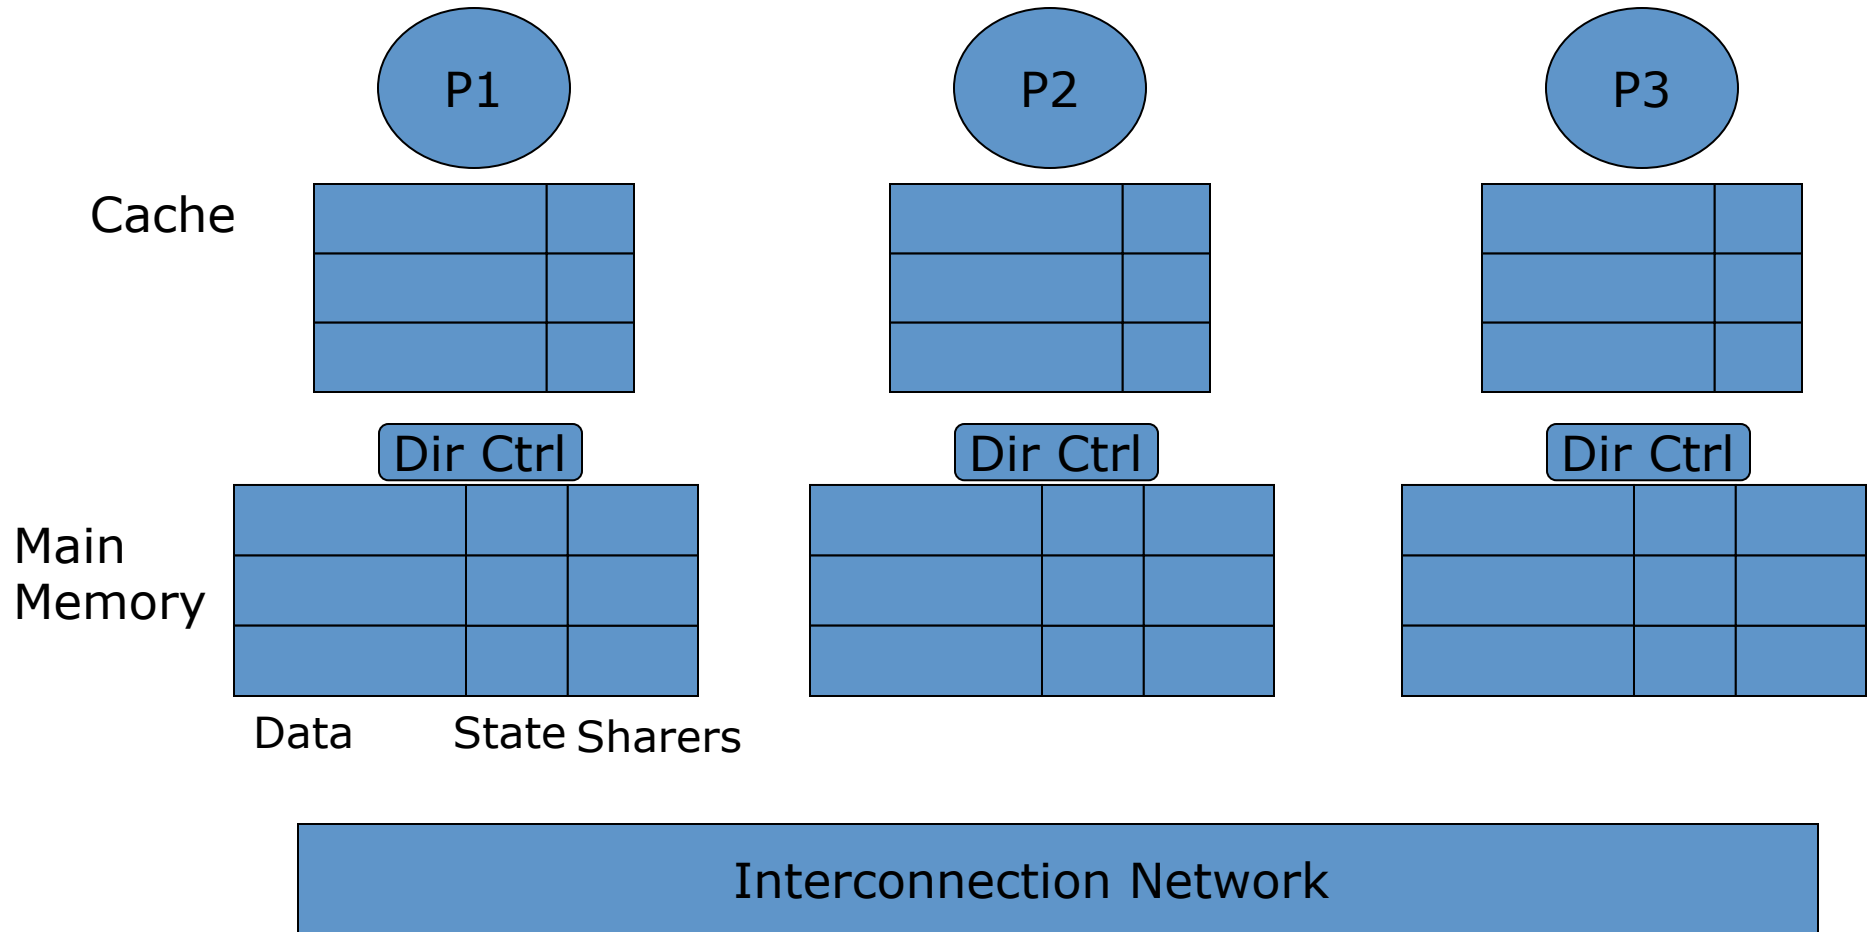
State Transition at the Directory



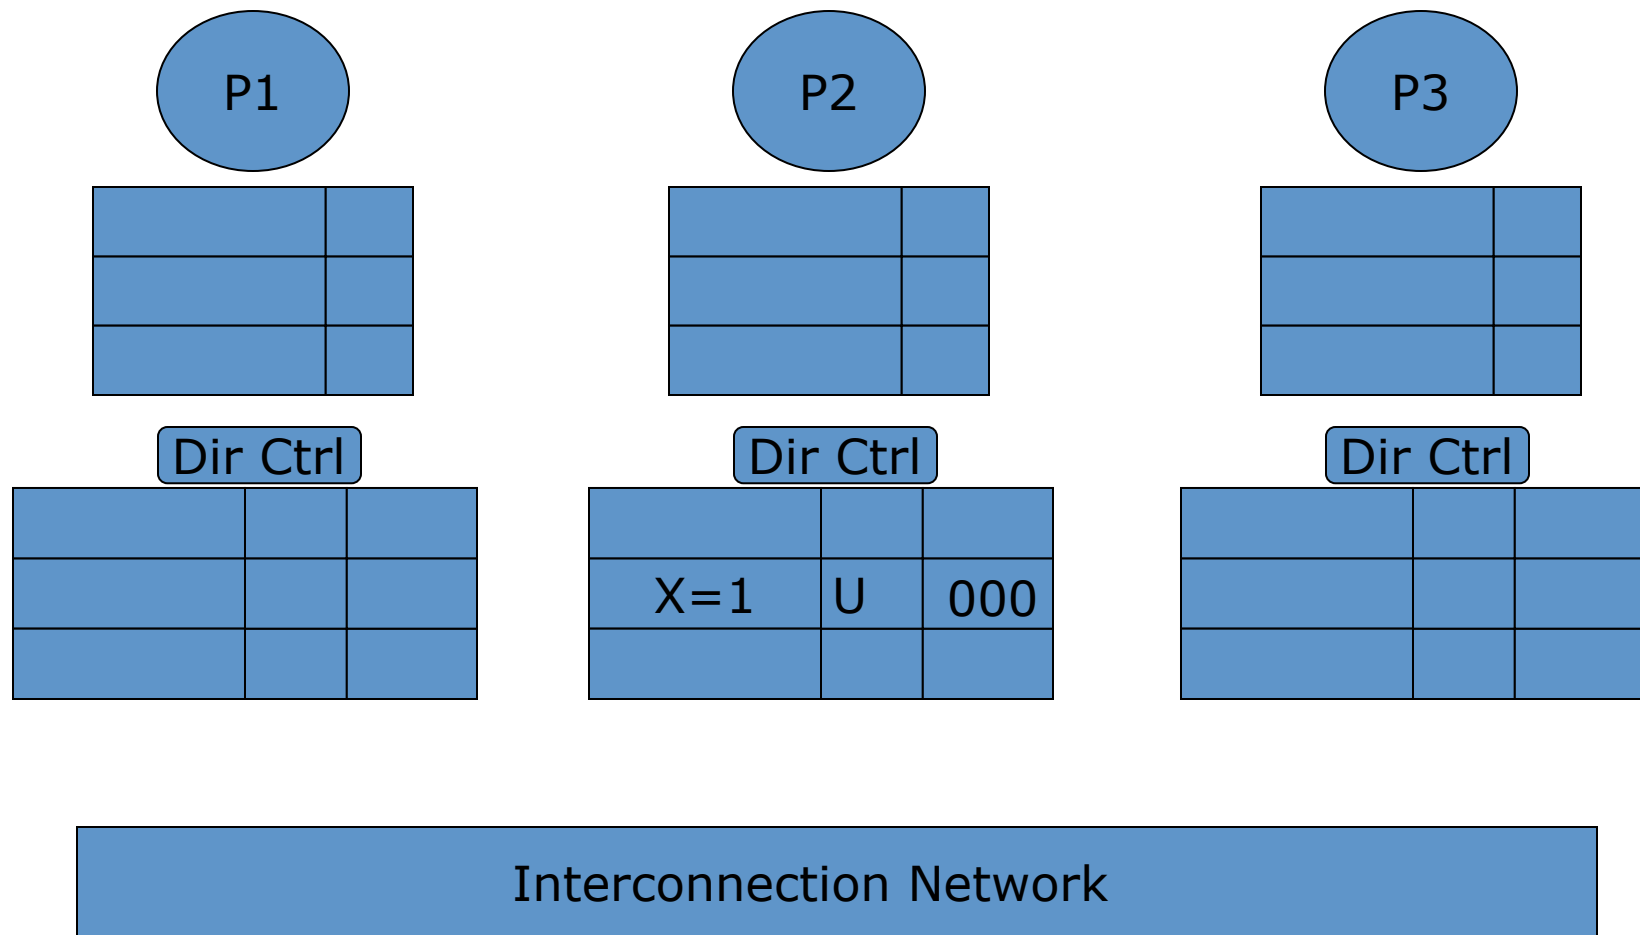
- Legend:
 - <Request-type>(<requestor>)/<Reply-type>(<destination>)
 - O = owner, A = requestor, Shr = sharers

Full Bit Vector Directory Protocol Animation

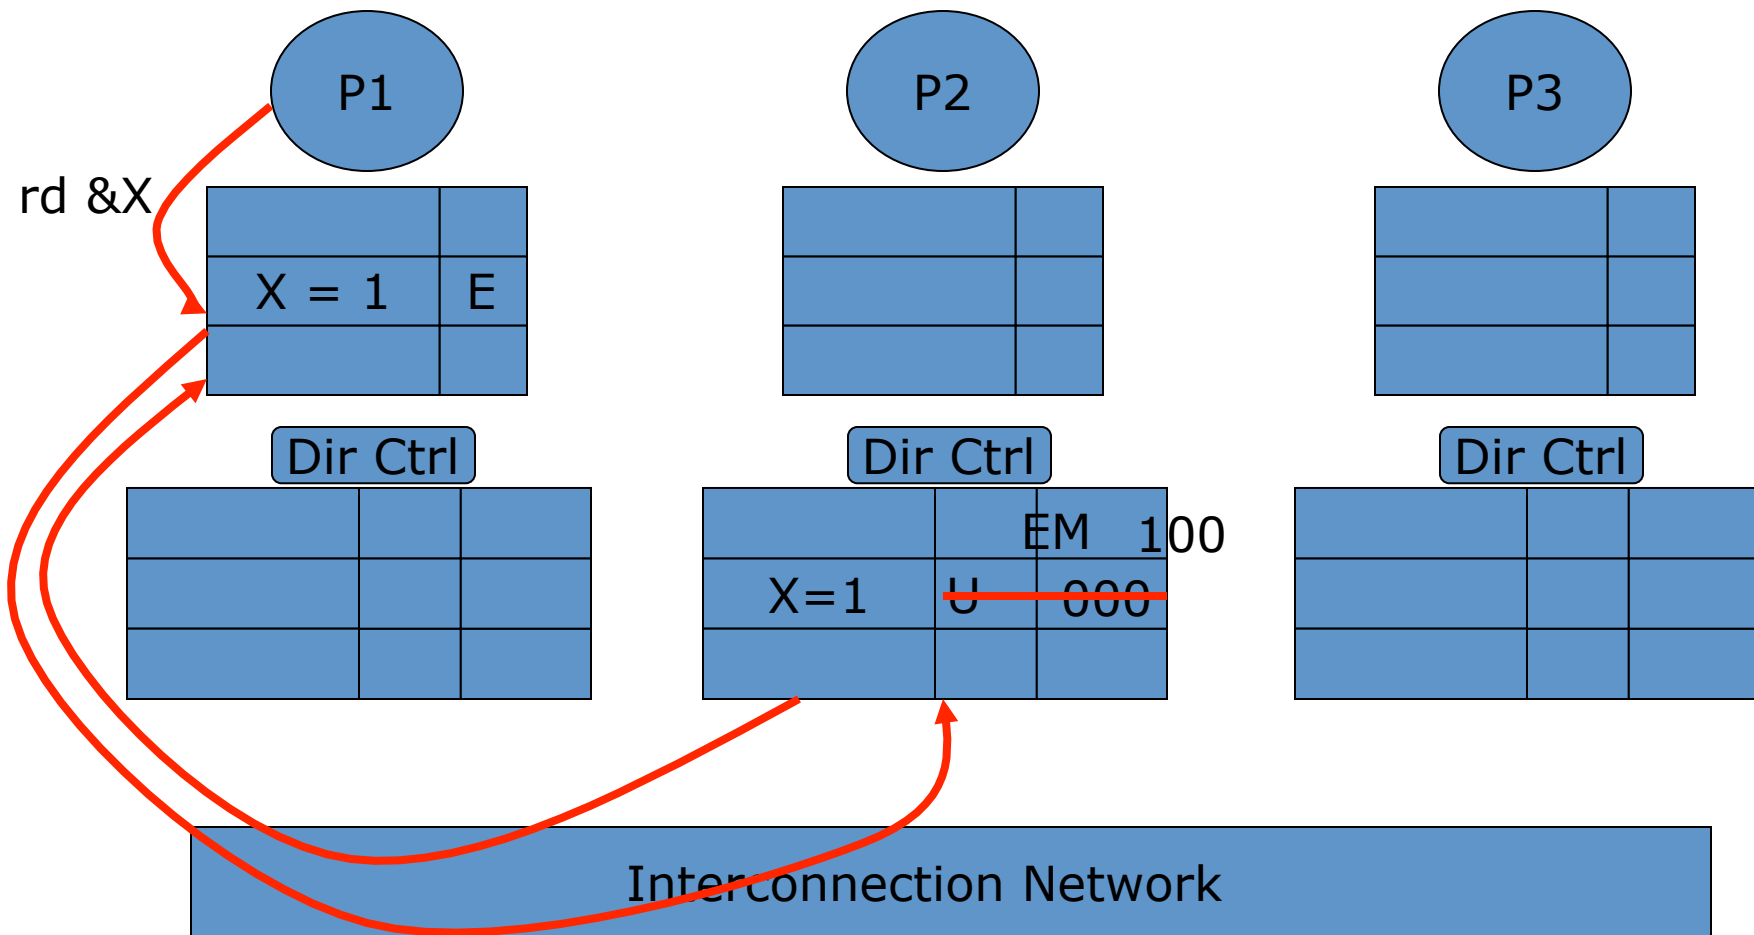
Full-bit Vector Protocol Visualization



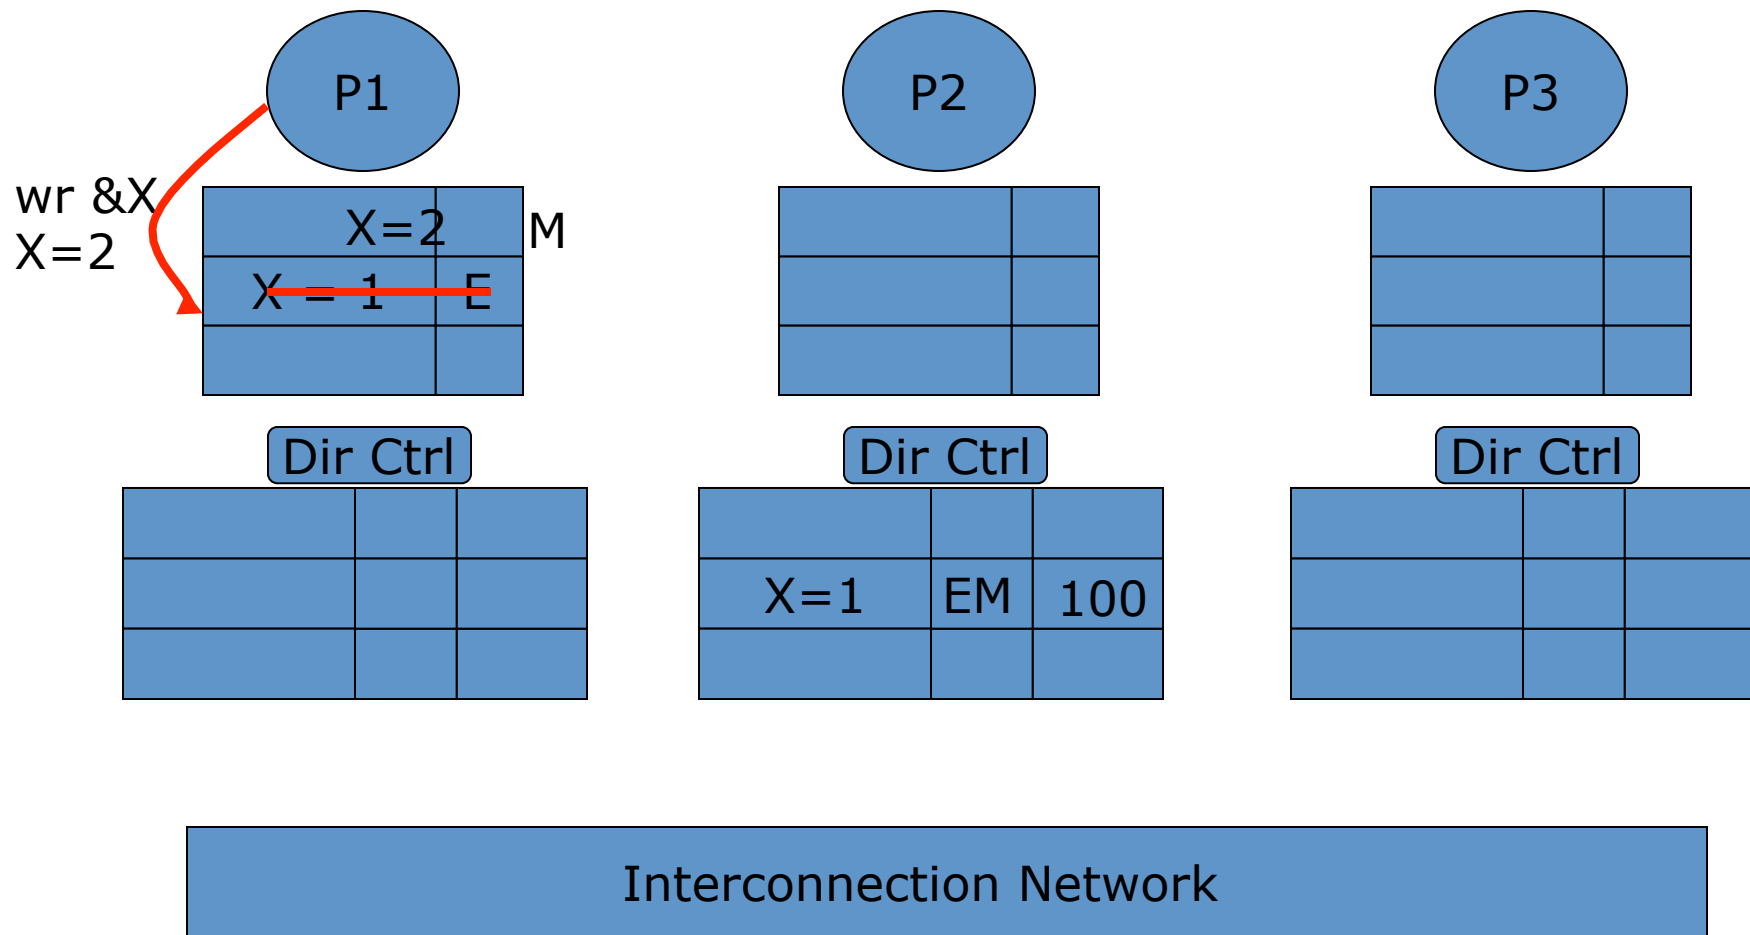
Full-bit Vector Protocol Visualization



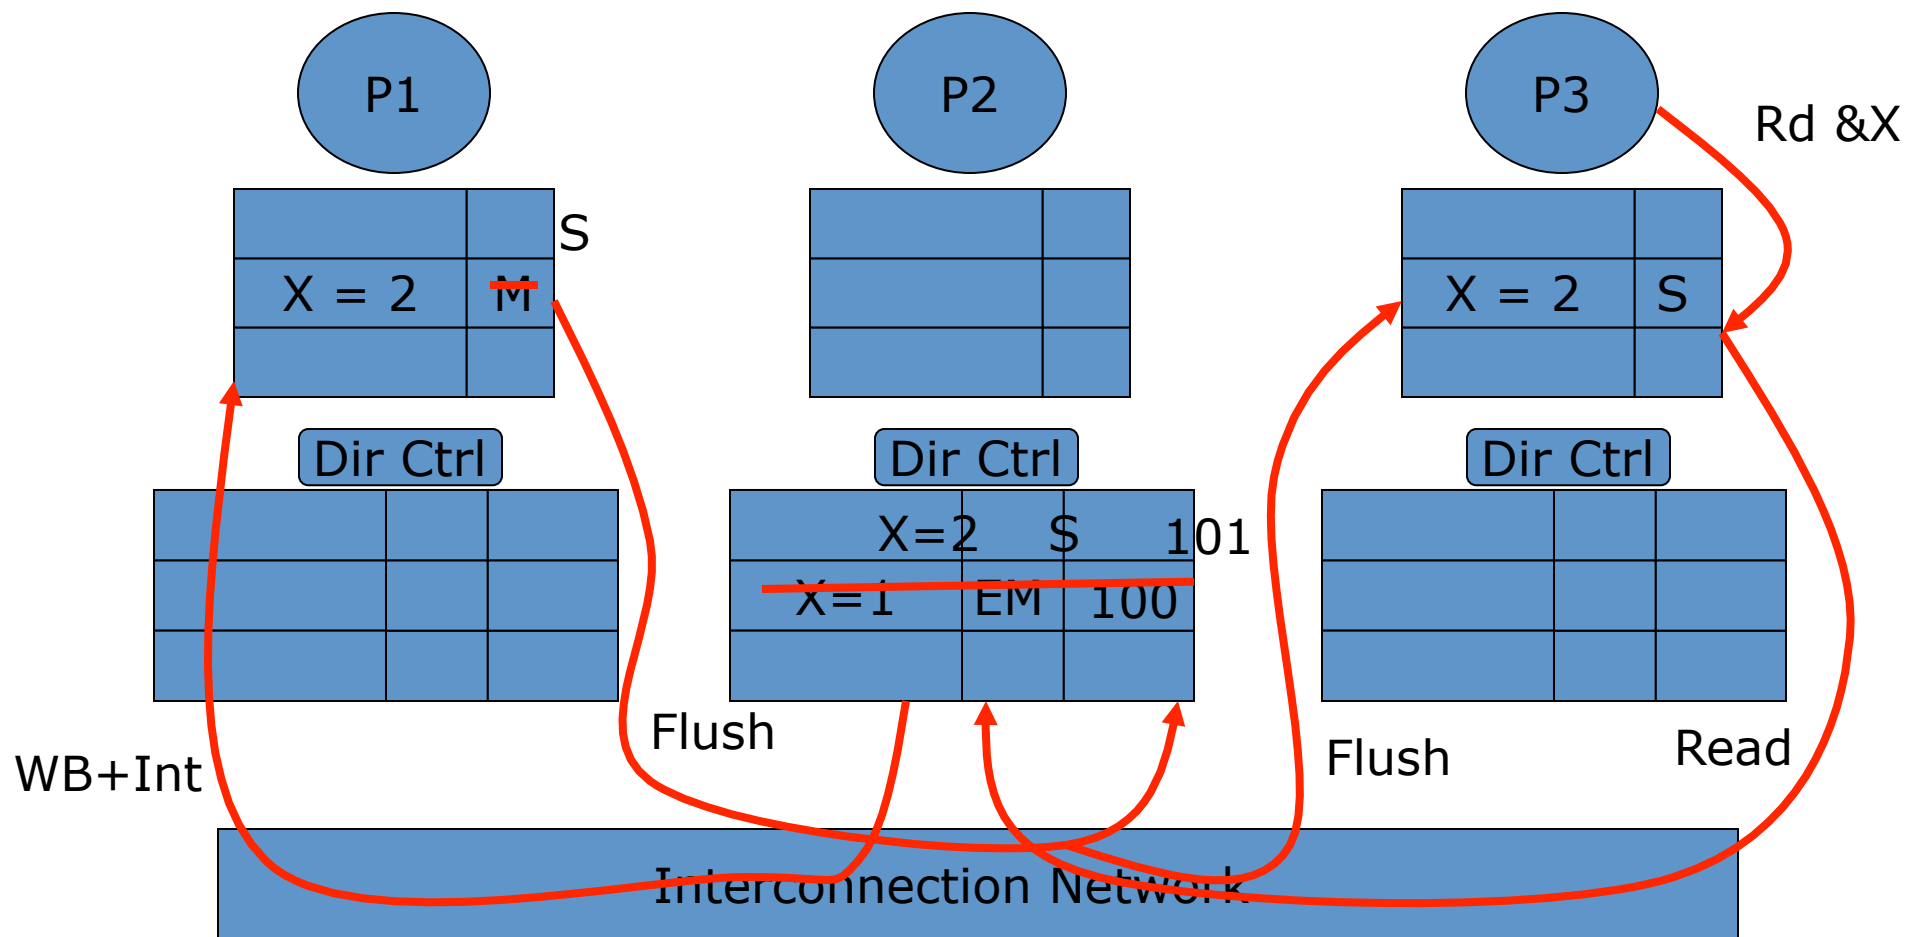
Full-bit Vector Protocol Visualization



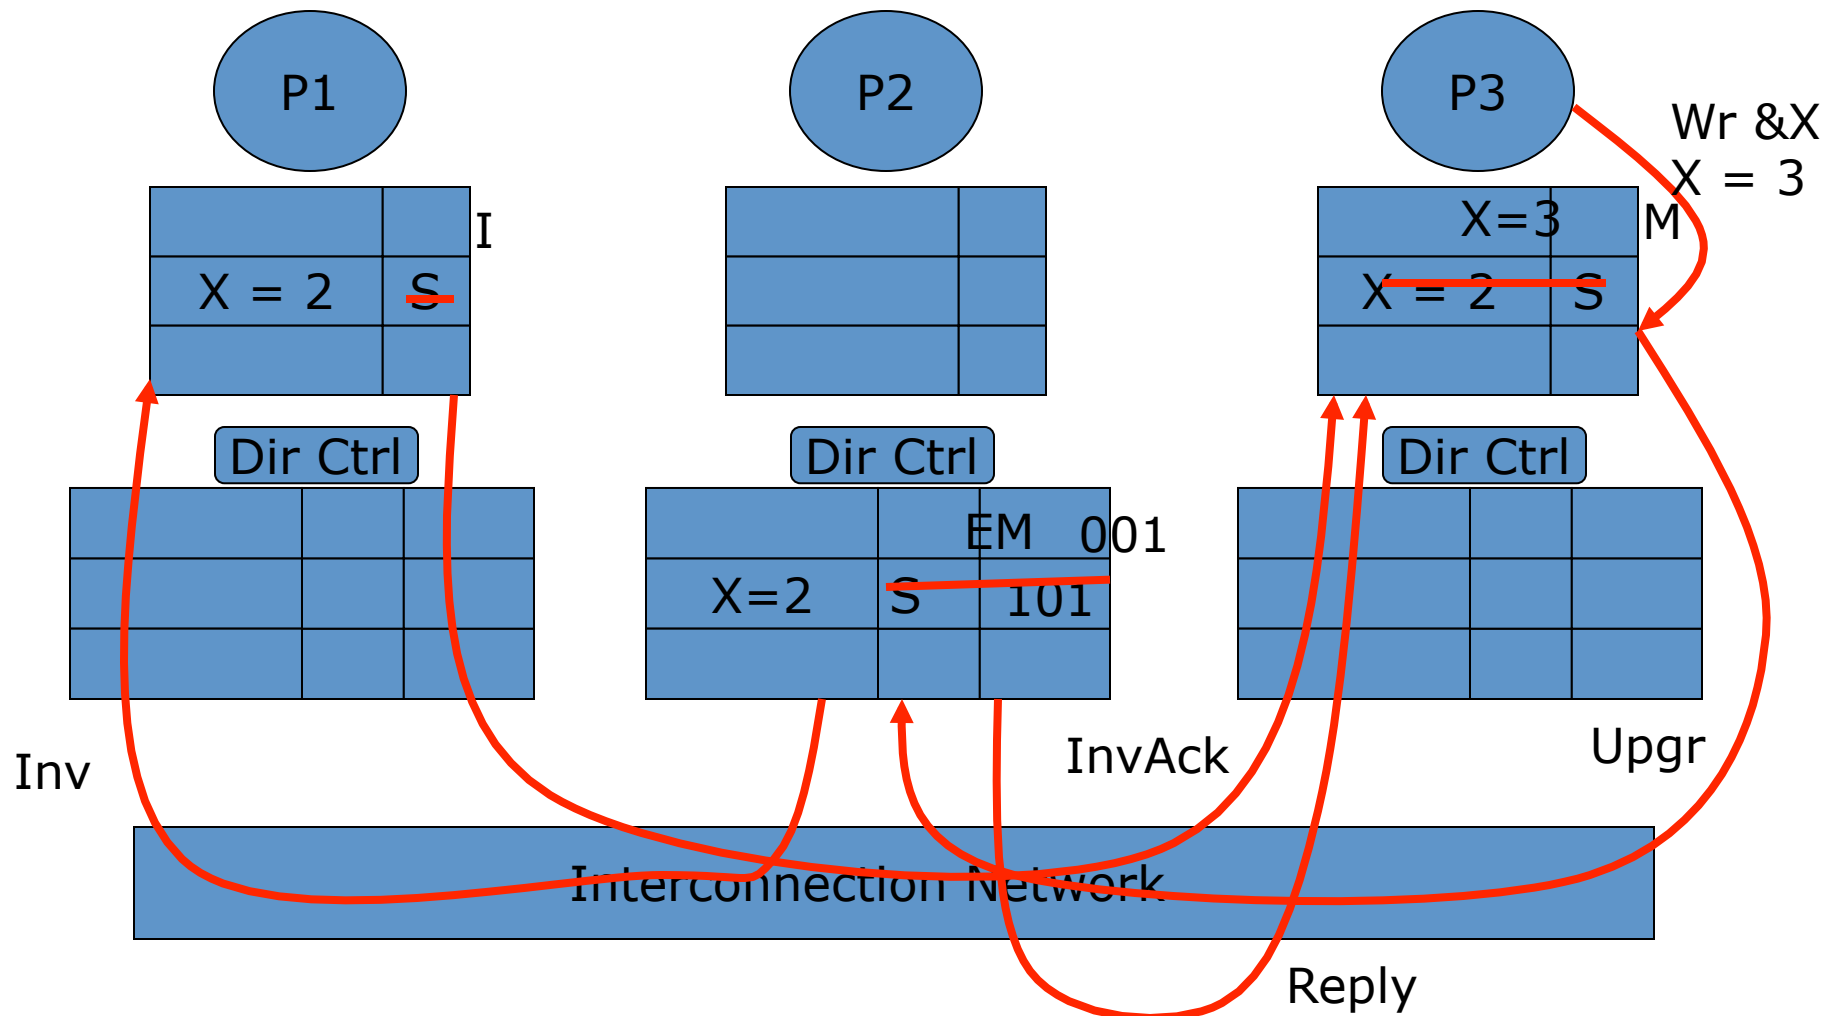
Full-bit Vector Protocol Visualization



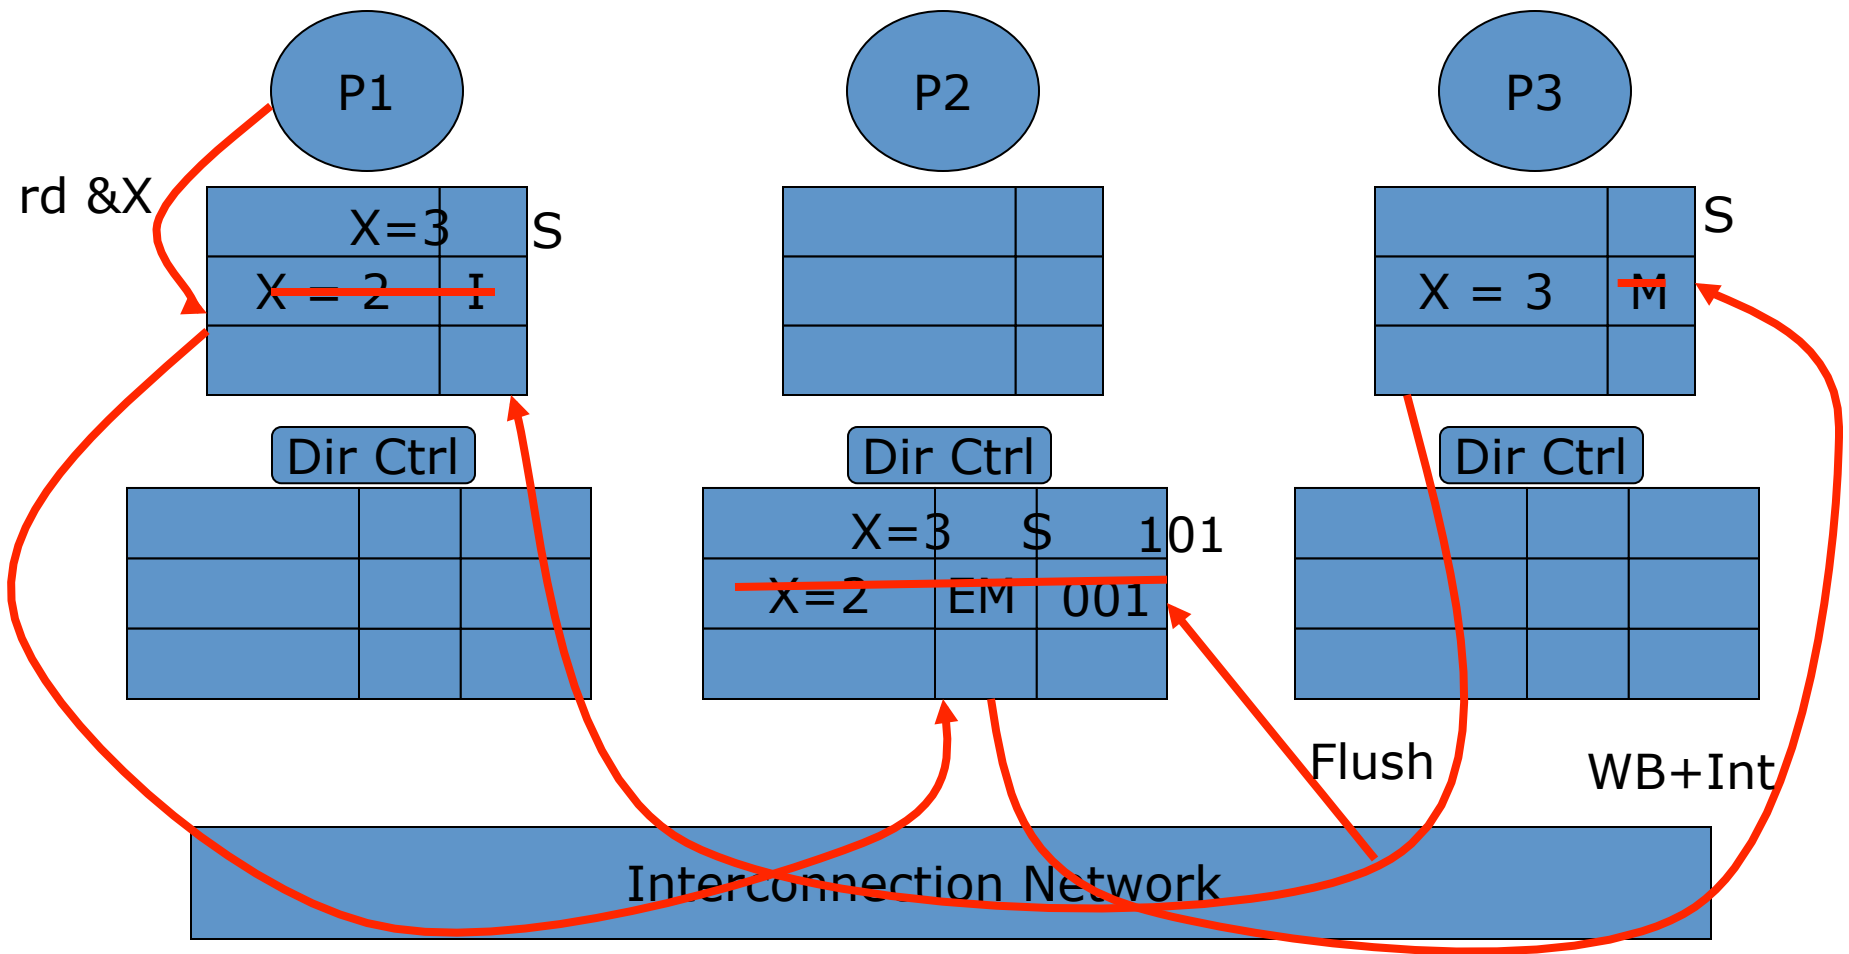
Full-bit Vector Protocol Visualization



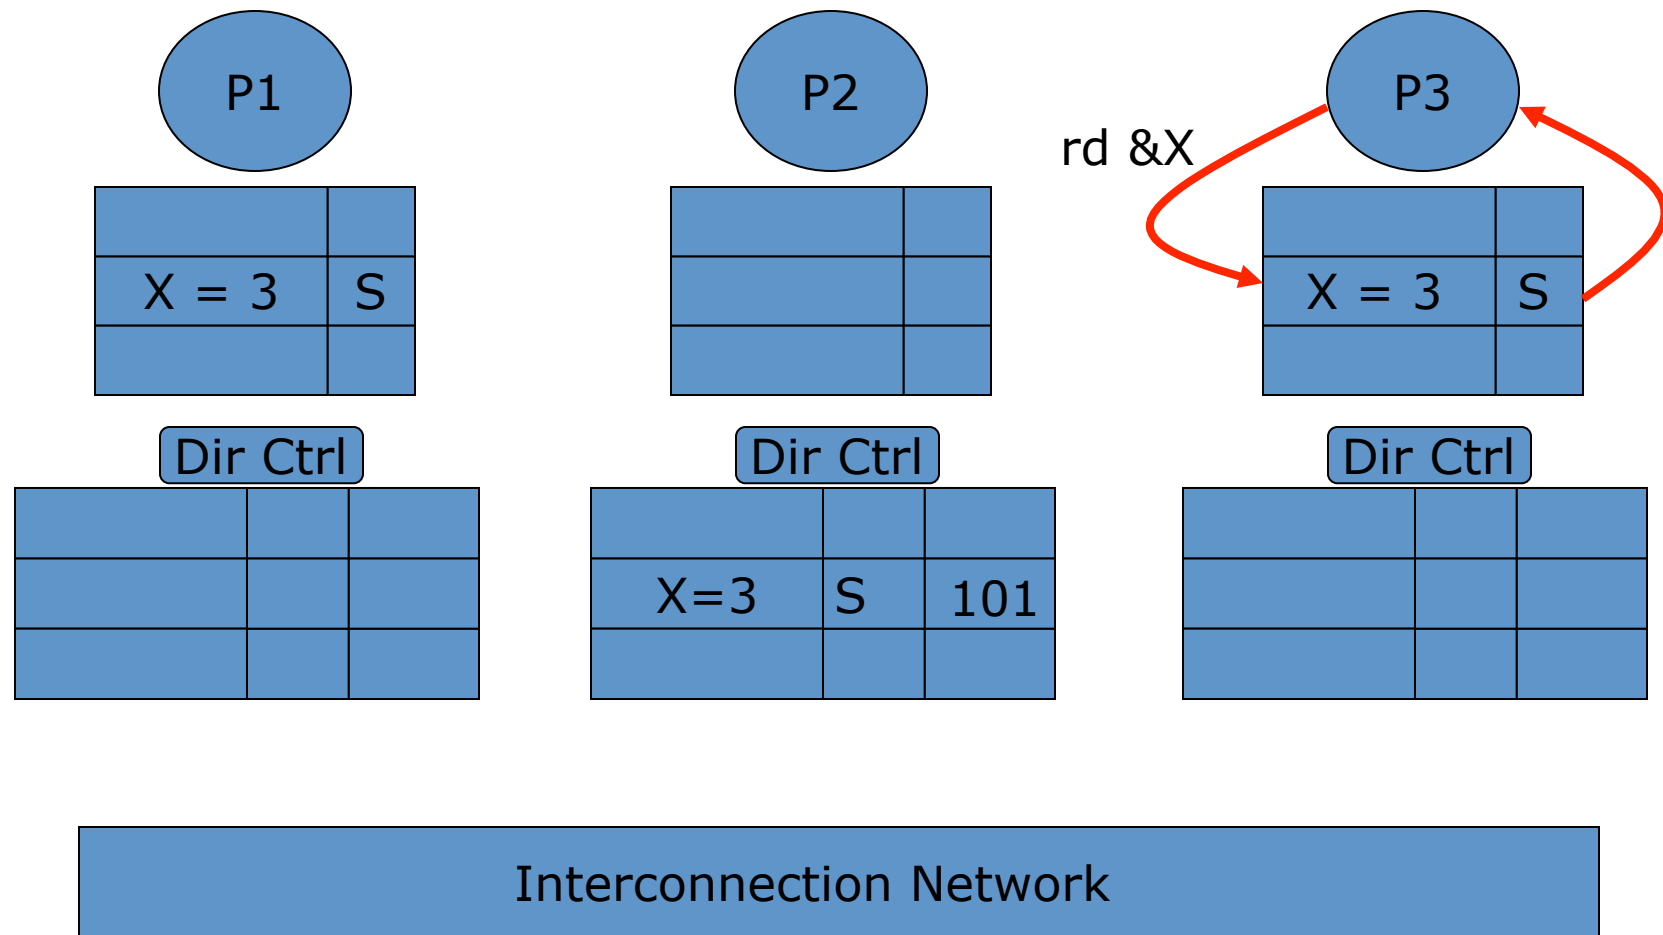
Full-bit Vector Protocol Visualization



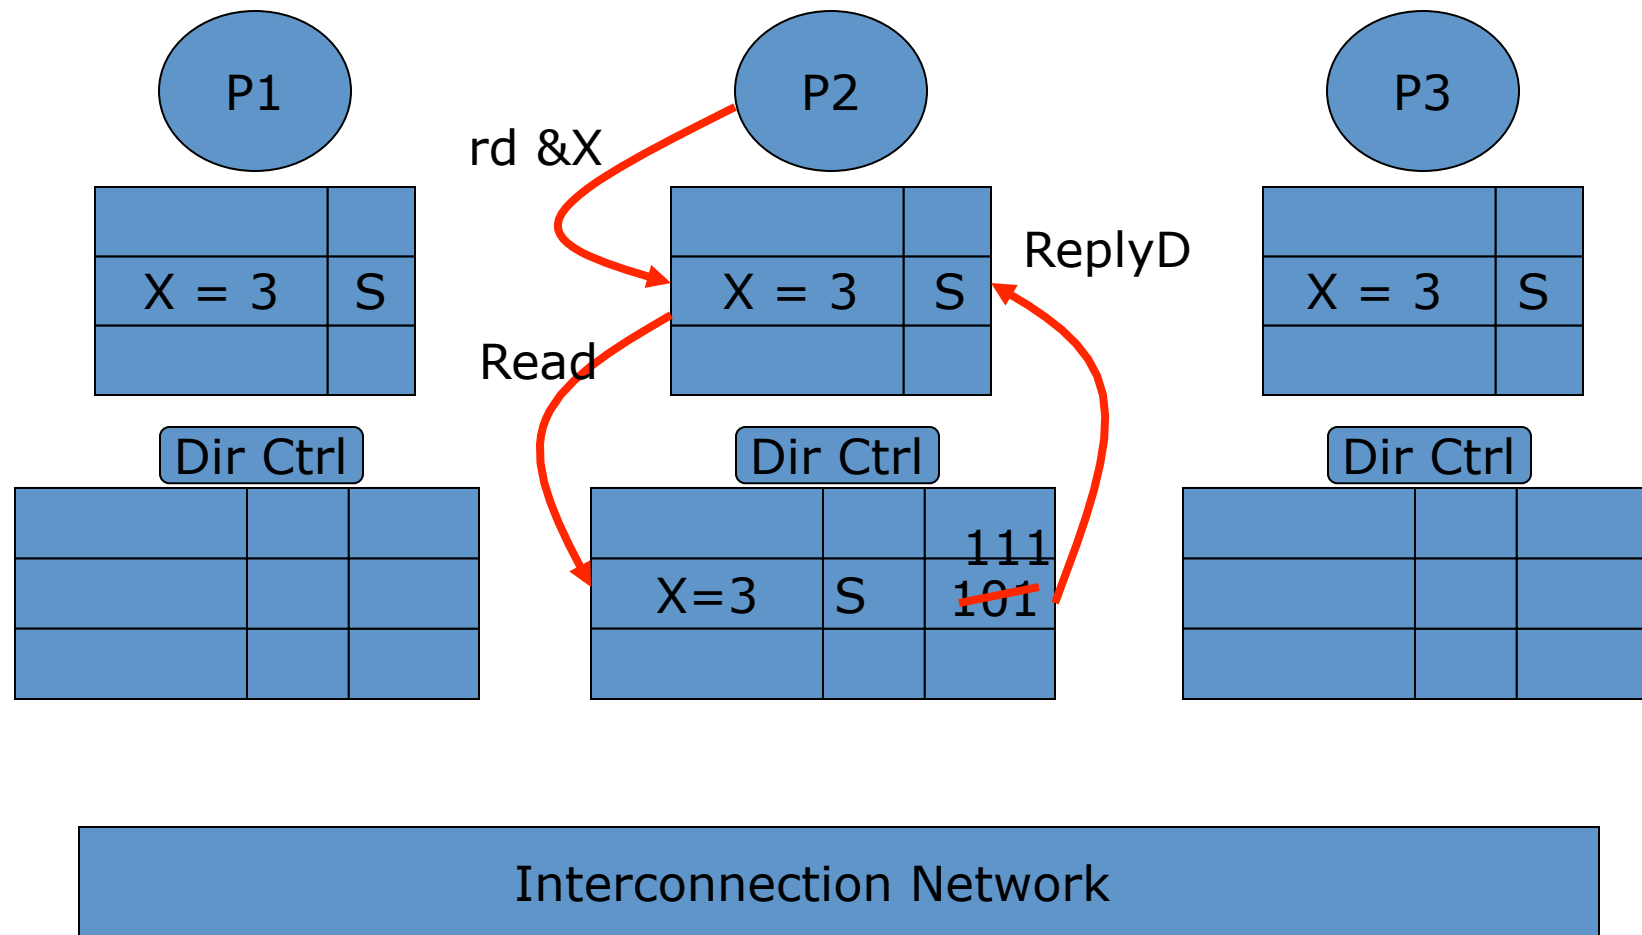
Full-bit Vector Protocol Visualization



Full-bit Vector Protocol Visualization



Full-bit Vector Protocol Visualization



Example

Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R1	E	-	-	EM, 100	Read (P1-> H), ReplyD (H->P1)	2
W1	M	-	-	EM, 100	-	0
R3	S	-	S	S, 101	Read (P3->H), WB+Int (H->P1), Flush (P1->H, P3)	3
W3	I	-	M	EM, 001	Upgr (P3->H), Reply (H->P3) // Inv (H->P1), InvAck(P1->P3)	3
R1	S	-	S	S, 101	Read (P1->H), WB+Int (H->P3), Flush (P1->H, P1)	3
R3	S	-	S	S, 101	-	0
R2	S	S	S	S, 111	Read(P2->H), ReplyD(H->P2)	2

Note

- Many transactions can be sent in parallel
- vs. snoopy coherence
 - Much more complicated (more msgs types)
 - Protocol races can occur (discussed later)
 - Some requests require 3 hops (vs. 2 in SMP)
 - FlushOpt in snoopy protocol MESI no longer needed because the block is supplied by home when clean
- Reply by home to node that issues a write request
 - Enables requestor to collect invalidation acknowledgements
 - Enables 3 hops instead of 4 hops
 - Called “Reply Forwarding” scheme

Overheads of Full Bit-Vector Scheme

- Traffic on a write
 - $O(\text{sharers})$
- Latency a write
 - $O(1)$ because Invalidations sent in parallel to sharers
 - But higher than SMP due to InvAck
- Storage overhead
 - doesn't scale well
 - E.g. 64-byte line implies
 - 64 nodes: 12.7% ovhd.
 - 256 nodes: 50% ovhd.; 1024 nodes: 200% ovhd.
 - Total storage overhead = $O(p^2)$
 - p bits per line, number of lines increases with p

Reducing Storage Overhead 1

- Simple Optimizations
 - Larger cache lines
- Coarse Vector
 - Use multiprocessor nodes (1 bit per group)
 - Storage overhead: $O(p^2/g)$ where g is group size
 - 256-procs, 4 per cluster, 128B line: 6.25% ovhd.
- Limited pointer scheme: storage $O(\text{numptr} * \log_2 p * p)$
 - Home has numptr pointers to sharers
 - $P=1024 \Rightarrow 10$ bit ptrs, even 100 pointers still save space
 - Few sharers mean few pointers enough (five or so)
 - When overflow (i.e. sharers $>$ pointers)?

Overflow Schemes for Limited Pointers

- Broadcast
 - Overflow bit set
 - Resort to broadcast
- Invalidate old sharers to fit max #sharers
 - on overflow, make room for new sharer by invalidating one of the old ones
- Coarse vector
 - On overflow, resort to coarse vector scheme
 - Write invalidation sent to all nodes that a bit corresponds to

Sparse Directories

- Observation: total number of cache entries \ll total amount of memory.
 - most directory entries have uncached (U) state
- Sparse directory: only keep directory entries that have non-U states
 - Organize directory as a cache
 - If a block will be cached, allocate an entry in the directory
 - However, it can overflow
 - Cached blocks can be replaced silently by caches
 - The directory does not see these replacement events, so it cannot free the entry
 - Handling overflow: send invalidations to all sharers when entry replaced (simpler)

Directory Information Format non Exclusive

- We can employ several formats simultaneously
- For example, in Origin2000 (2 procs form a node):
 - (1) Exclusively cached: dir entry is pointer to that specific processor
 - (2) Shared: bit vector, each bit stores node id (not processor id)
 - Invalidation to a Hub is broadcast to both processors in the node
 - Two sizes and storage
 - 16-bit format (32 procs), kept in main memory DRAM
 - 64-bit format (128 procs), extra bits kept in extension memory
 - (3) Larger machines: coarse vector, each bit corresponds to p/64 nodes
 - Invalidation is sent to all Hubs in that group, each hub bcast to its 2 procs
- Machine chooses between bit vector and coarse vector dynamically

Physical Storage Organization for Directory

- Refer to the discussion at the end of Section 11.2.1

Module 11.2 – Handling Protocol Races

Assumptions So Far

- We have assumed
 - Directory state reflects the most up to date state of caches
 - Messages due to a request are processed atomically
- In reality, one of or both conditions may be violated
 - => **protocol races** can occur
 - Some protocol races can be handled in a simple way
 - Others are trickier to handle
- We will discuss how protocol races can be handled
 - Purpose of discussion: illustrating approaches for dealing with protocol races
 - Discussing all possible races is not the goal (I.e., some races are left out in the discussion)

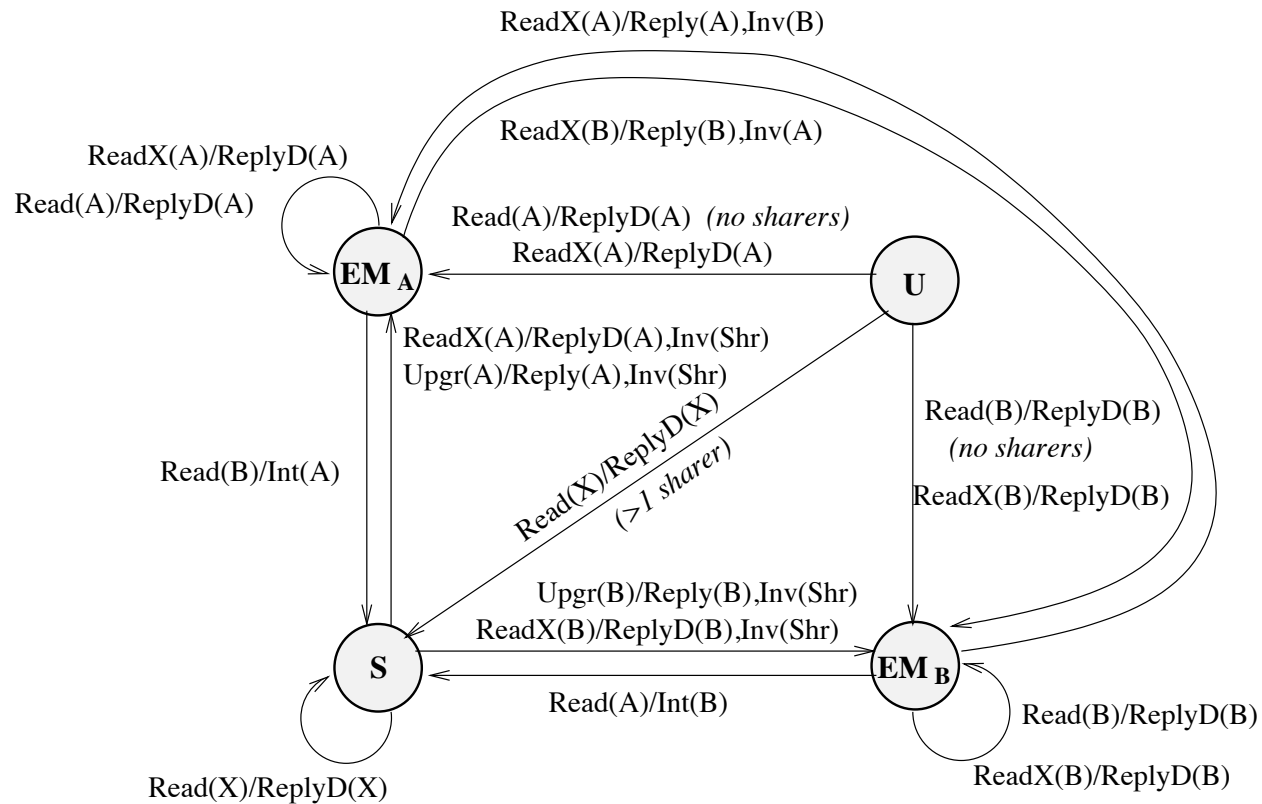
Handling Races: Out-of-sync Directory

- Home sends invalidation to a node that has replaced the block silently
 - The node can reply with InvAck
- Home receives a read request from a node that is already a sharer from the home point of view
 - Directory can reply with data (ReplyD)
- Home receives a read/write request from a node that the home thinks as the owner (EM state)
 - If the block was clean, it has been replaced silently
 - If the block was dirty, the Flush (or write back) has yet to reach the home
 - What should the home do?
 - Wait? The block may never come
 - Reply with data? The data may be stale

Handling Races: Out-of-sync Directory


- The home alone cannot solve this. Individual nodes must participate in the solution
 - Keep outstanding transaction buffer for any Flush message
 - Require Home to acknowledge the receipt of a Flush
 - Delay Read/ReadX request to a block that is still being written back
- Hence, home only receives Read/ReadX to a block that is not being written back
 - It can send ReplyD

Protocol Modification



- Split EM into two states: EM_A and EM_B
- Distinguish handling Read/ReadX requests from current owner or others

Handling Races: Non-Atomic Messages




- 1: A sends a read request to Home
- 2: Home replies with data (but the message gets delayed)
- 3: B sends a write request to Home
- 4: Home sends invalidation to A, and it arrives before the ReplyD
- Called “Early Invalidation” race

How Should A Respond?


- Two incorrect ways to respond:
 - A replies with InvAck,
 - B thinks that its write propagation is complete
 - A receives a ReplyD and places the block in its cache (the block that should have been invalidated)
 - A ignores the Invalidation message
 - The message is lost, and write propagation has failed to occur
- Solution
 - Brute force: avoids overlapped handling of requests. Home waits until it receives Ack from all parties (Home-centric)
 - Allows overlapping but ask nodes to participate (Requestor-assisted)
 - Node keeps outstanding transaction buffer
 - It does not entertain requests (to the same block) until the current transaction is completed

Processing a Read Request

Read to clean block:



Read to Excl/Modified block:



Legend:

- ┌- - -┐ processing period (home-centric)
- ┌- - -┐ processing period (requestor-assisted)

- Home-centric (Read to clean block)
 - Directory enters a transient state
 - Home replies with data
 - Requestor receives data, sends Ack to Home
 - Home closes transaction (transitions to a stable state, update sharing vector)
 - Cons: too much serialization at home, transaction closed late, and it requires Ack

Requestor-assisted (Read to clean block)

- Directory sends ReplyD, then closes transaction
- Requestor buffers/NACKs all new requests until the current Read transaction is completed (ReplyD received)


Figure 10.9: The processing period of a Read request.

Processing a Read Request


- Home-Centric (Read to EM block)
 - Requestor sends Read to Home
 - Home enters a transient state, sends intervention to Owner
 - Owner flushes block to Home and Requestor
 - Requestor sends Ack back to Home
 - Home closes transaction (transitions to Shared state, updates sharing vector)
- Requestor-assisted (Read to EM block)
 - Requestor sends Read to Home
 - Home enters a transient state, sends intervention to Owner
 - Home cannot close the transaction yet, because in the final state (Shared), it must have a clean copy of the block
 - Owner flushes block to Home and Requestor
 - Upon receiving the block from Owner, Home closes transaction

Processing a Write Request


ReadX to Unowned block:




ReadX to Shared block:



ReadX to Excl/Modified block (no WB race):



ReadX to Excl/Modified block (WB race):



Legend:

- processing period (home-centric)
- processing period (requestor-assisted)

- Home-centric (ReadX to U block)
 - Requestor sends ReadX request
 - Home replies with data
 - Requestor sends Ack
 - Home closes transaction
- Requestor-assisted (ReadX to U blk)
 - Requestor sends ReadX request
 - Home sends ReplyD and closes transaction
- Home-centric (ReadX to Shared blk)
 - Home enters transient state and sends Inv messages
 - InvAcks must be collected at Requestor and Requestor notifies Home, or collected at Home
 - Then, home closes transaction
- Requestor-assisted (ReadX to Shared blk)
 - Home sends Invs and closes the transaction
 - InvAcks collected at Requestor
 - Requestor buffers/NACKs new requests until it receives all InvAcks

Figure 10.10: The processing period of a ReadX request.

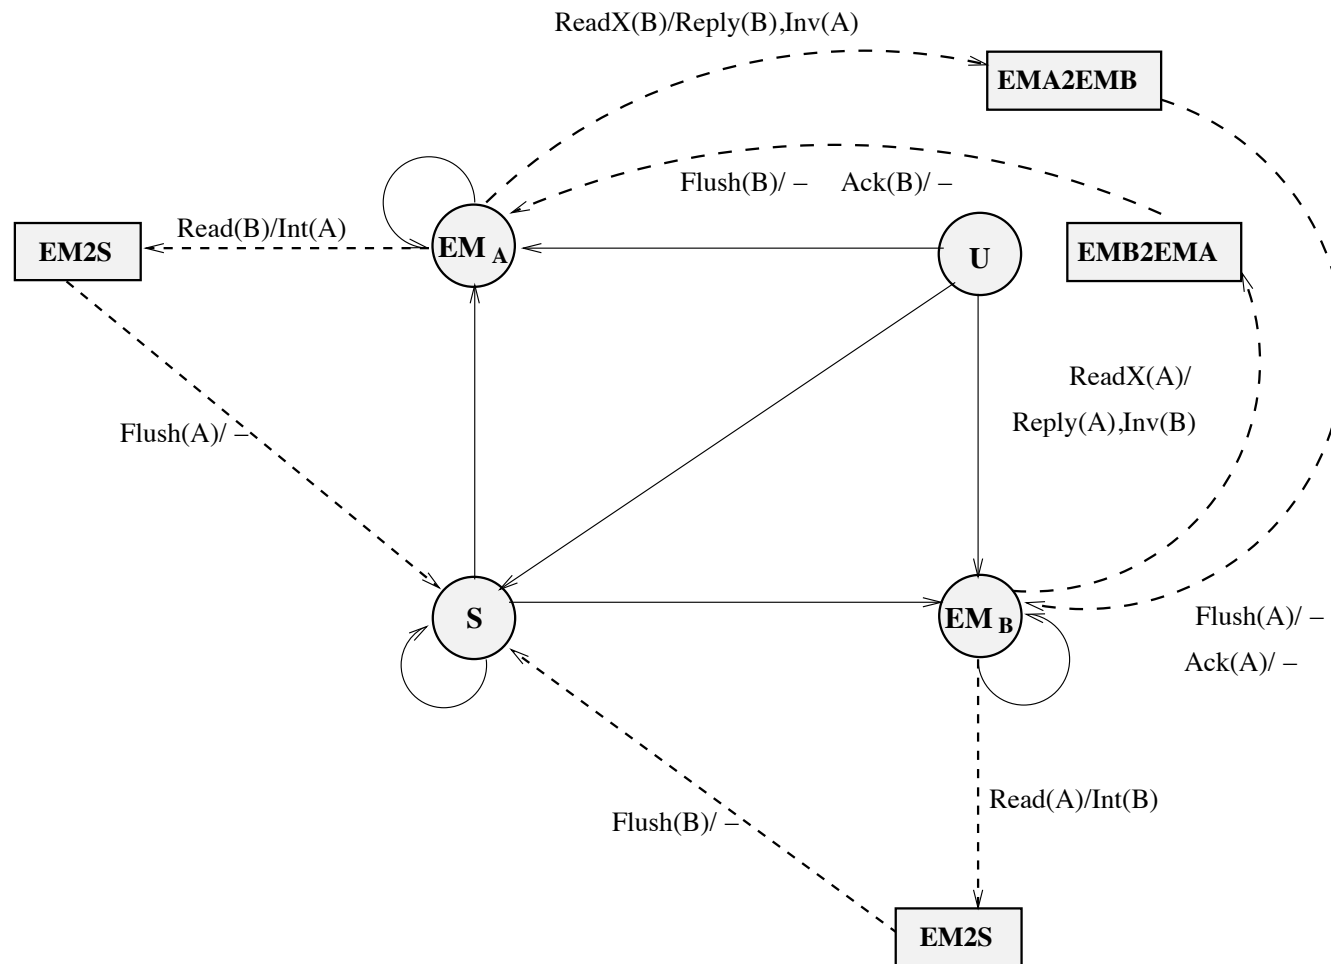
Processing a Write Request

- Home-Centric (ReadX to EM block)
 - Home enters transient state and sends Inv messages
 - InvAcks must be collected at Requestor and Requestor notifies Home, or collected at Home
 - Then, home closes transaction
- Requestor-Assisted (ReadX to EM block)
 - Requestor sends ReadX request to Home
 - Home sends Invalidation request to Owner, and closes the transaction
 - Owner flushes the block to Requestor
 - Requestor buffers/NACKs new requests until it receives the block from the old Owner

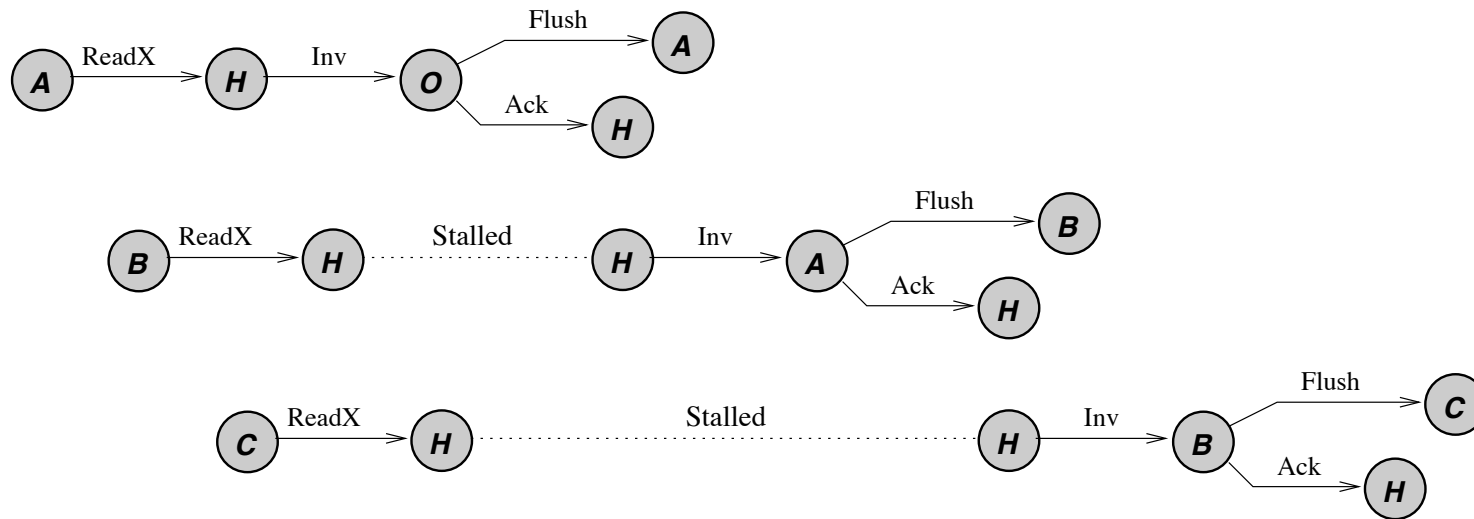
Processing a Write Request

- However, a problem remains:
 - What if the current Owner no longer has the block?
 - Either it had it in Modified state and it is writing the block back
 - or it had it in Exclusive state and it has evicted the block silently
 - Home cannot close the transaction yet, as it may have to supply the block
 - Hence, it can close the transaction late, after it receives Ack from Owner
- Transient states needed each time a transaction cannot be closed instantly
 - EM2S: transitioning from EM to S
 - EM2EM: transitioning from EM to EM (but different owner)

New Coherence State Diagram



Write Propagation and Serialization



- Write propagation is achieved through invalidation
- Multiple writes to a block are serialized by the protocol
 - Transaction closes after the Ack from current Owner is received by Home
 - New ReadX request is not served until the previous ReadX request is closed
 - This provides write serialization

Synchronization Support

- Atomic instructions
 - Can be supported by entering a transient state at Home until the instruction is completed
- LL/SC pair
 - No change is needed at the directory
- In-memory atomic operation
 - If directory controller can perform simple operations (increment, compare, add), it can perform atomic operation there, and
 - Allocate synchronization variables in a non-cacheable page
 - Pros: no invalidation and subsequent cache misses occur ($O(p)$ complexity)
 - Cons: performs worse under little-contention scenarios (due to a long distance and lack of temporal locality)

Memory Consistency Models

- SC implementation
 - Write completion detected when all InvAcks are collected
 - Prefetching and load speculation are applicable
- As # processors grow
 - Average latency of a cache miss increases
 - Harder to hide it
 - Hence, benefits of more relaxed consistency models increase vs. SC

Module 11.3 – Cache-Based Directory Protocol

Flat, Cache-based Schemes

- How they work:
 - Sharers are linked as a linked list
 - home only stores the head pointer
 - cache stores the next sharer in the list
 - on read, add yourself to head of the list (comm. needed)
 - on write, propagate chain of invalids down the list
- Scalable Coherent Interface (SCI) IEEE Standard
 - doubly linked list

Assumptions

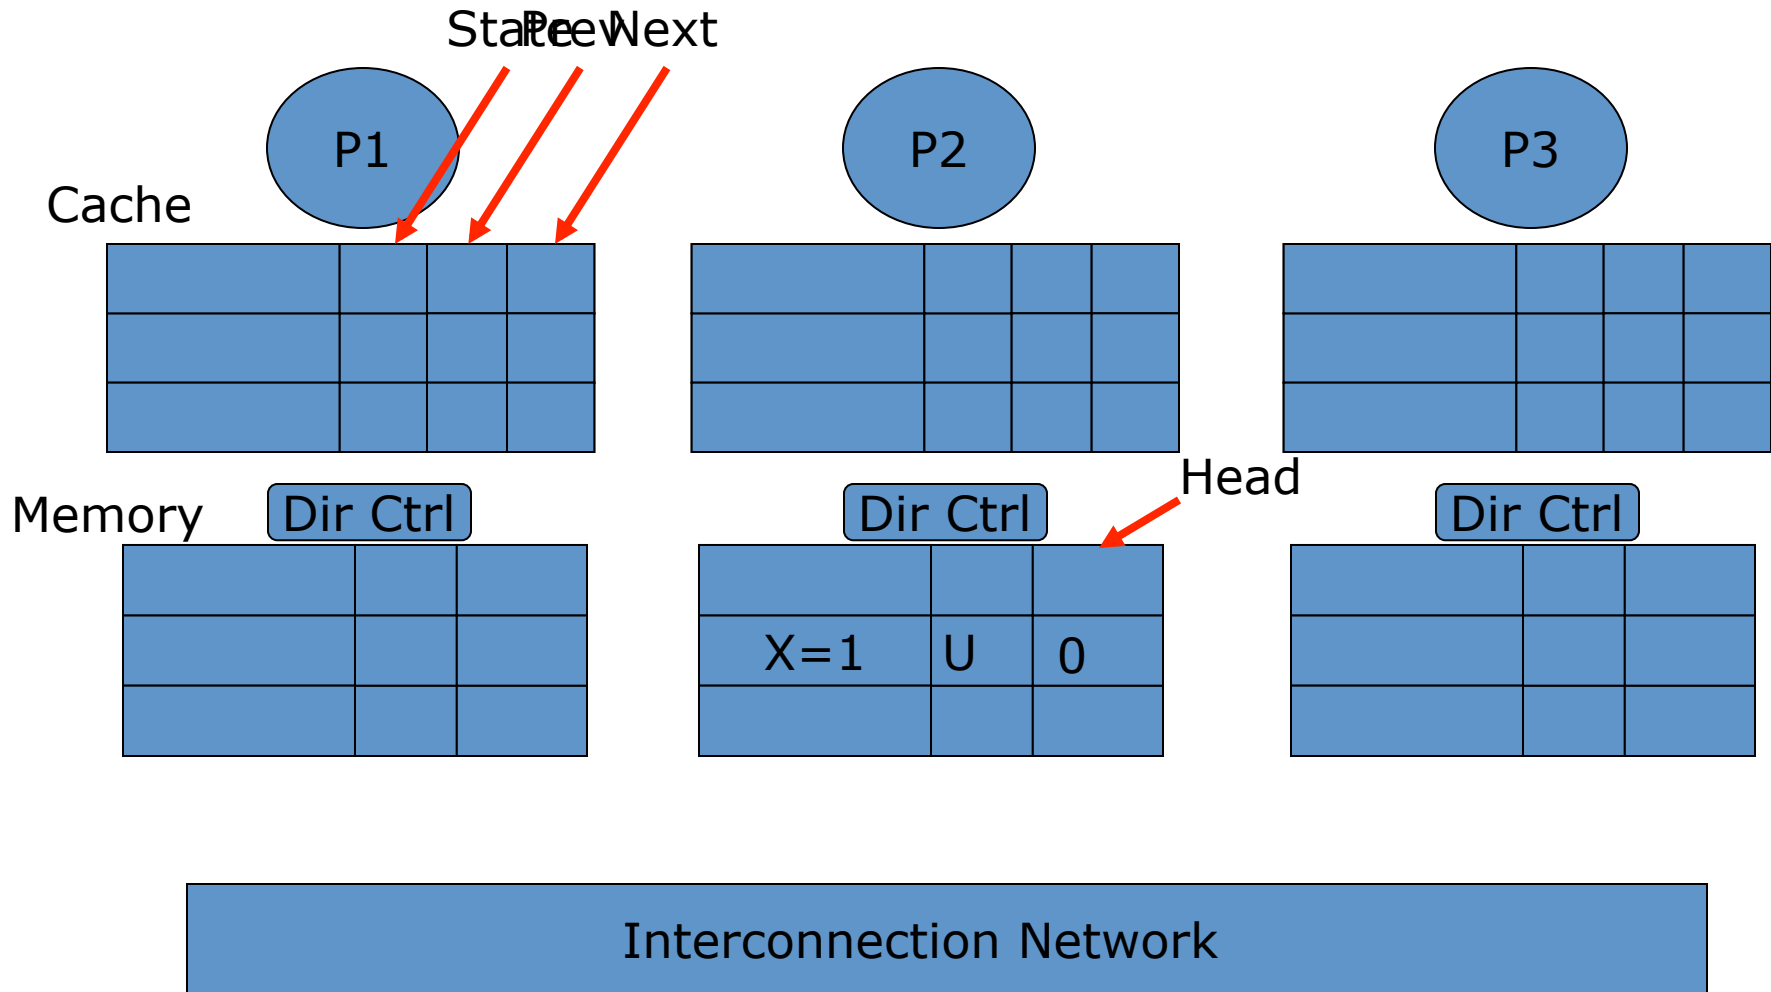
- SCI (Scalable Coherence Interface) protocol
 - IEEE standard, ratified in 1993
 - 3 sets of protocols: minimal, typical, full
 - 7 state bits, 29 stable states + many pending states
- For illustration: Simple SCI (SSCI)
 - Retain similarity with full-bit vector protocol:
 - MESI states in the cache
 - U,S,EM states in the memory directory
 - Replace the presence bits with a pointer
 - Similar features with SCI
 - Overall protocol operation
 - Doubly linked list
 - Many possible race conditions
 - Mostly ignored in the illustration

SSCI: Example

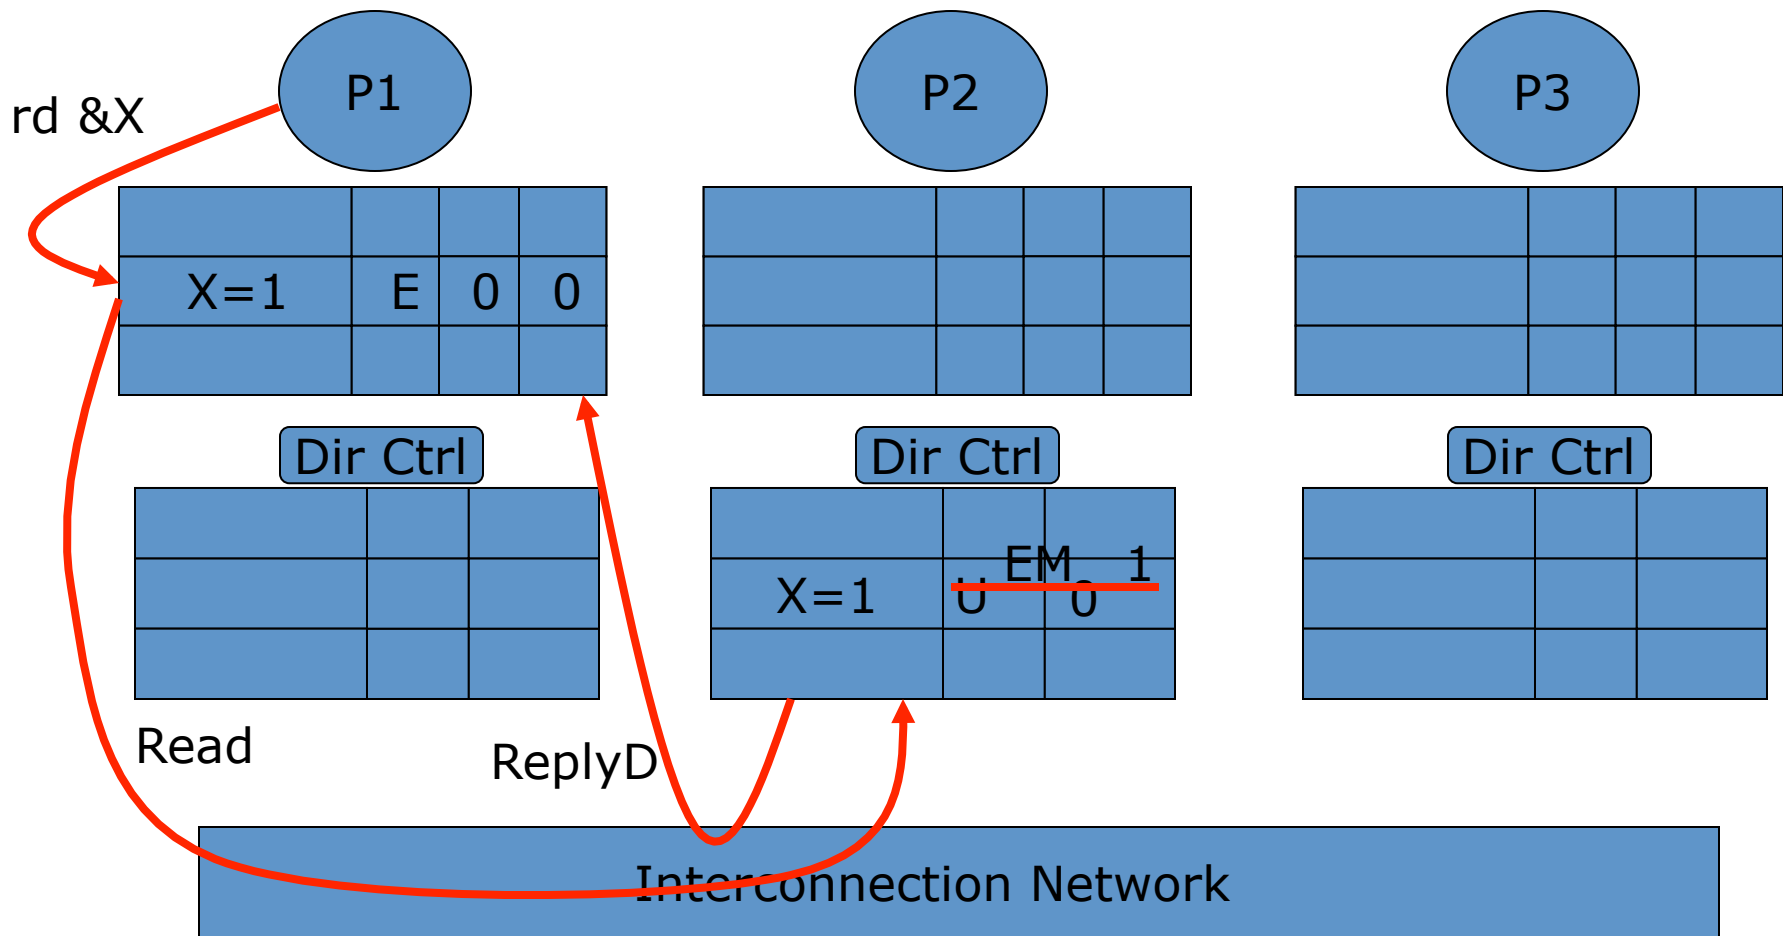
- Coherence Network Transactions: add
 - WB+Int+UpdPtr
 - UpdPtr: update next/prev/head pointers

SSCI Protocol Animation

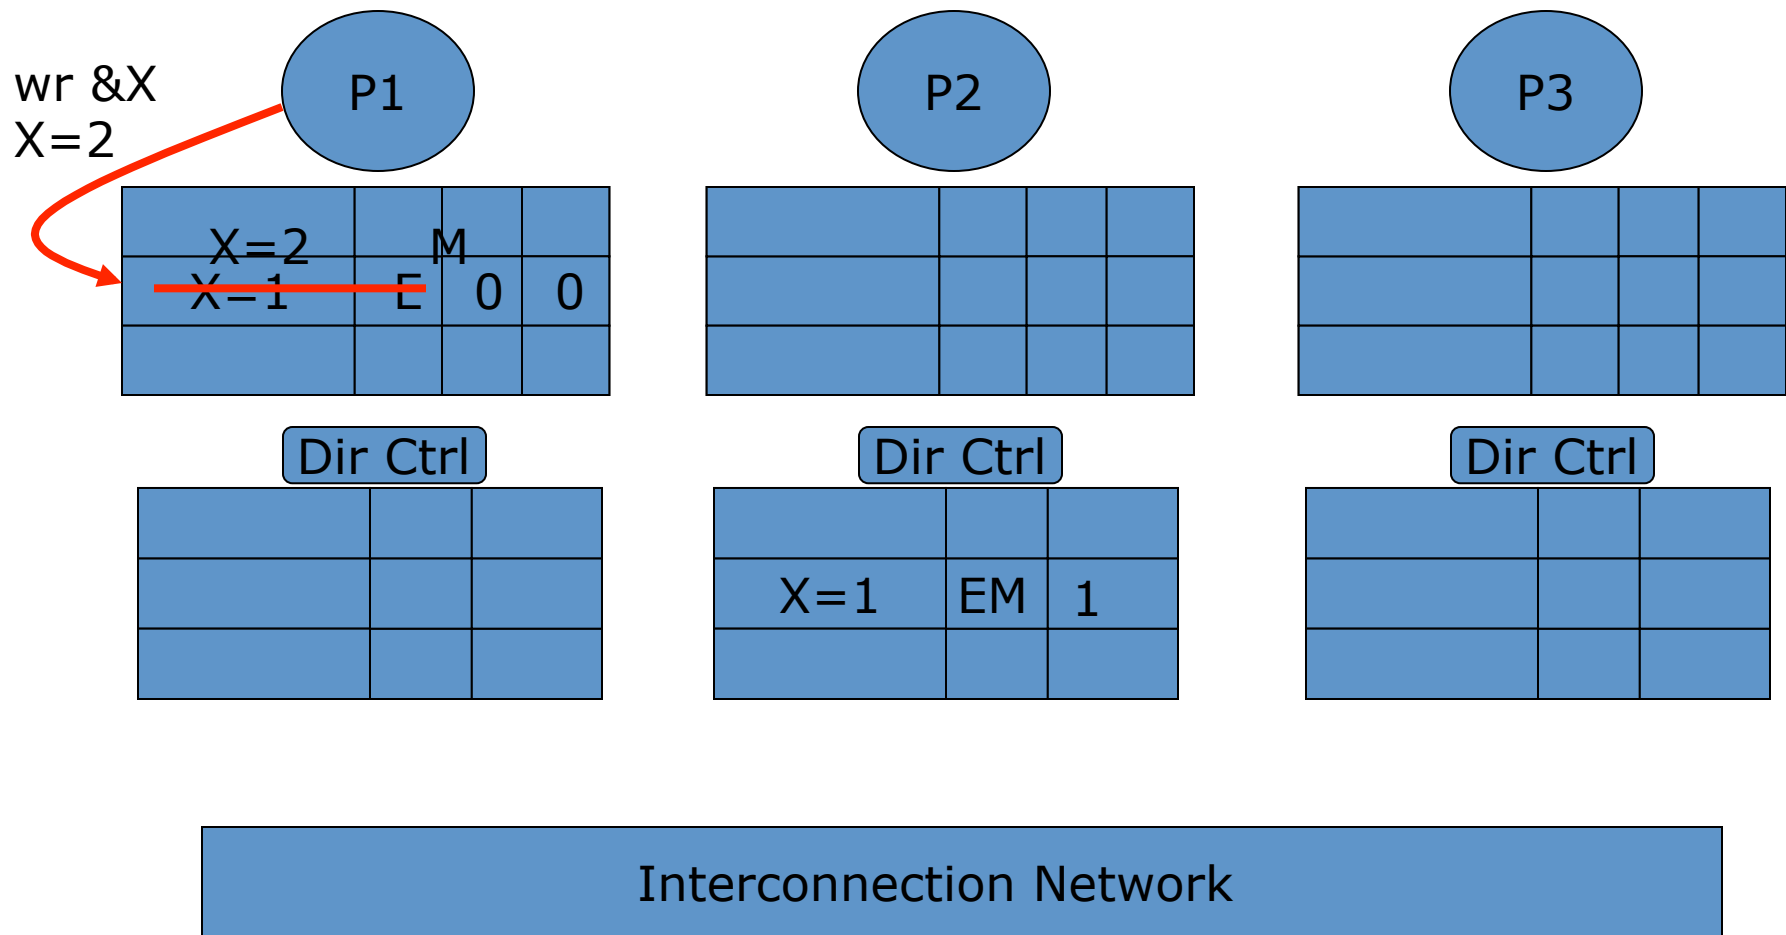
SSCI Protocol Visualization



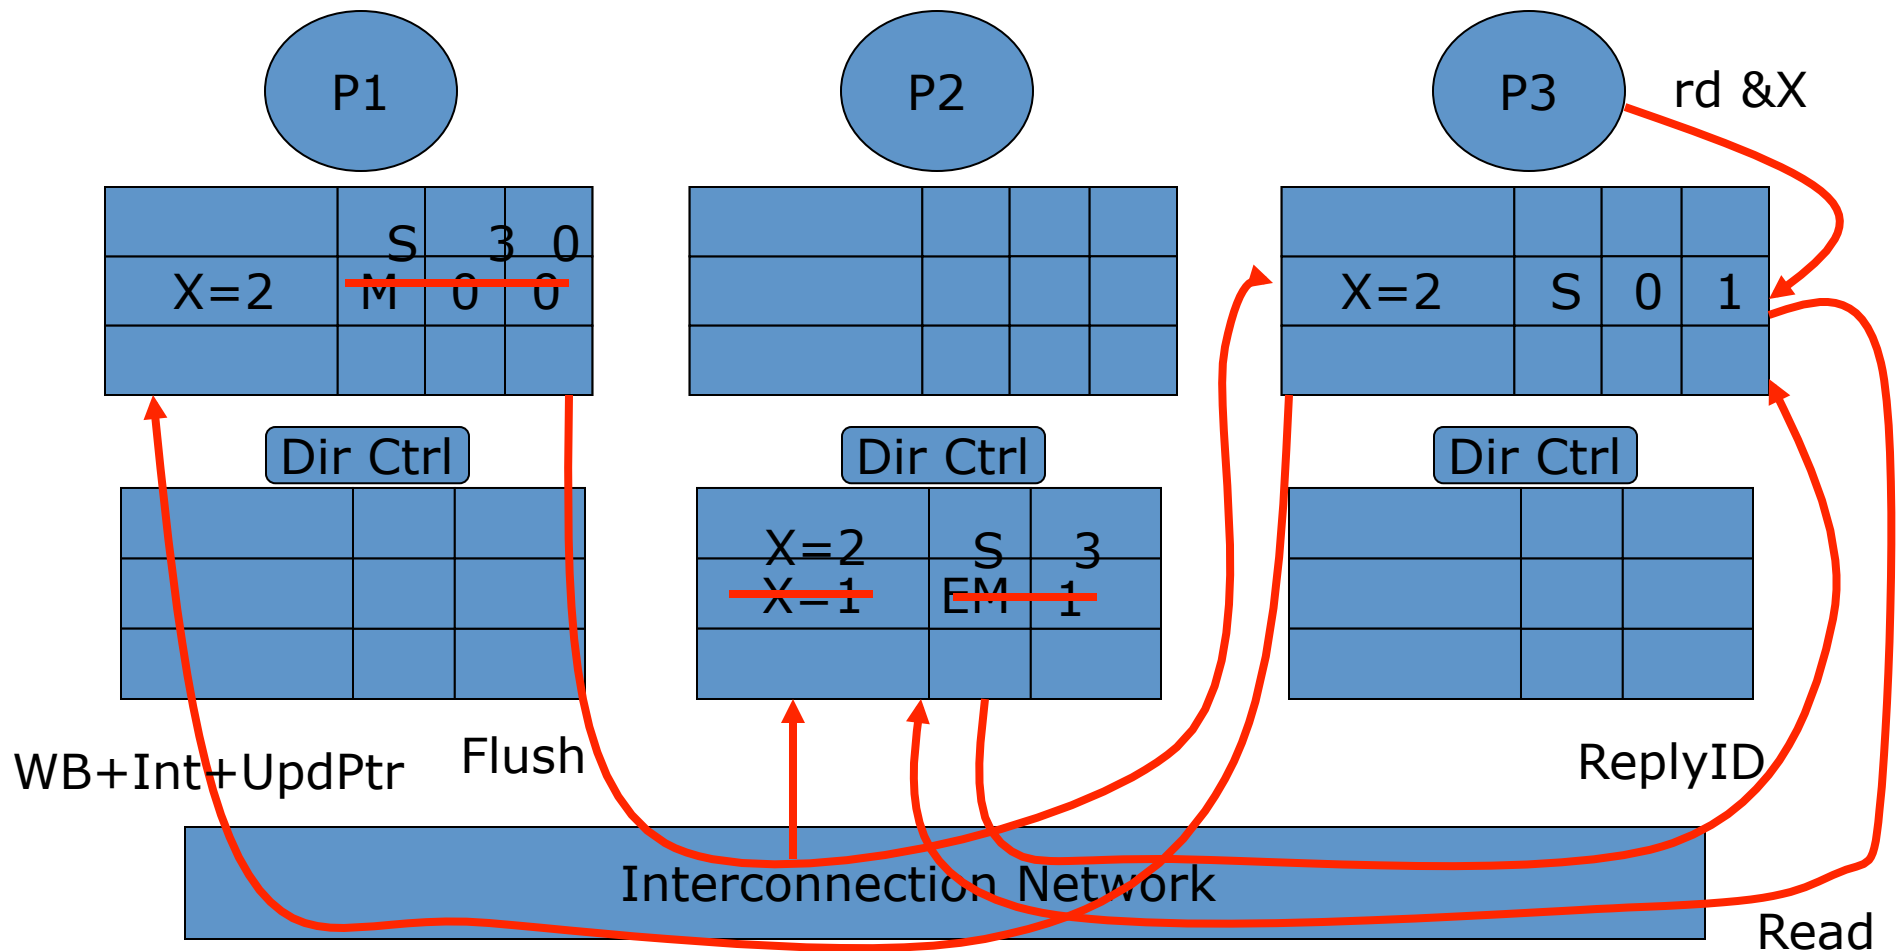
SSCI Protocol Visualization



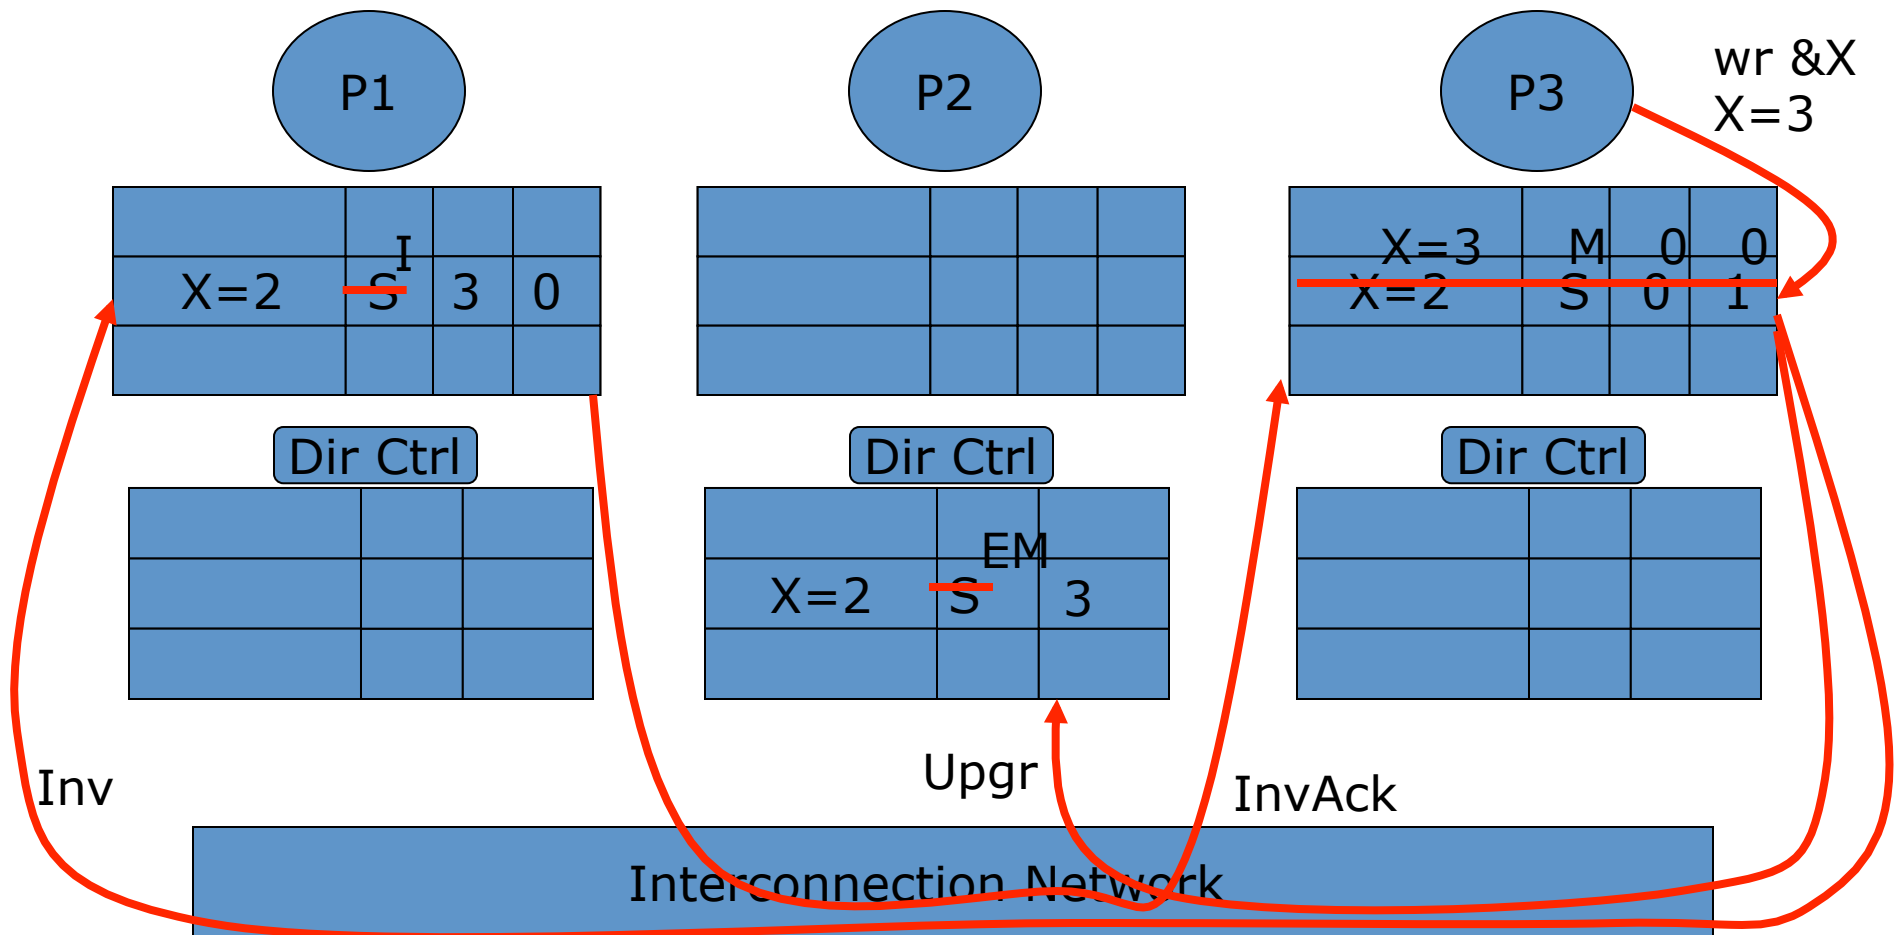
SSCI Protocol Visualization



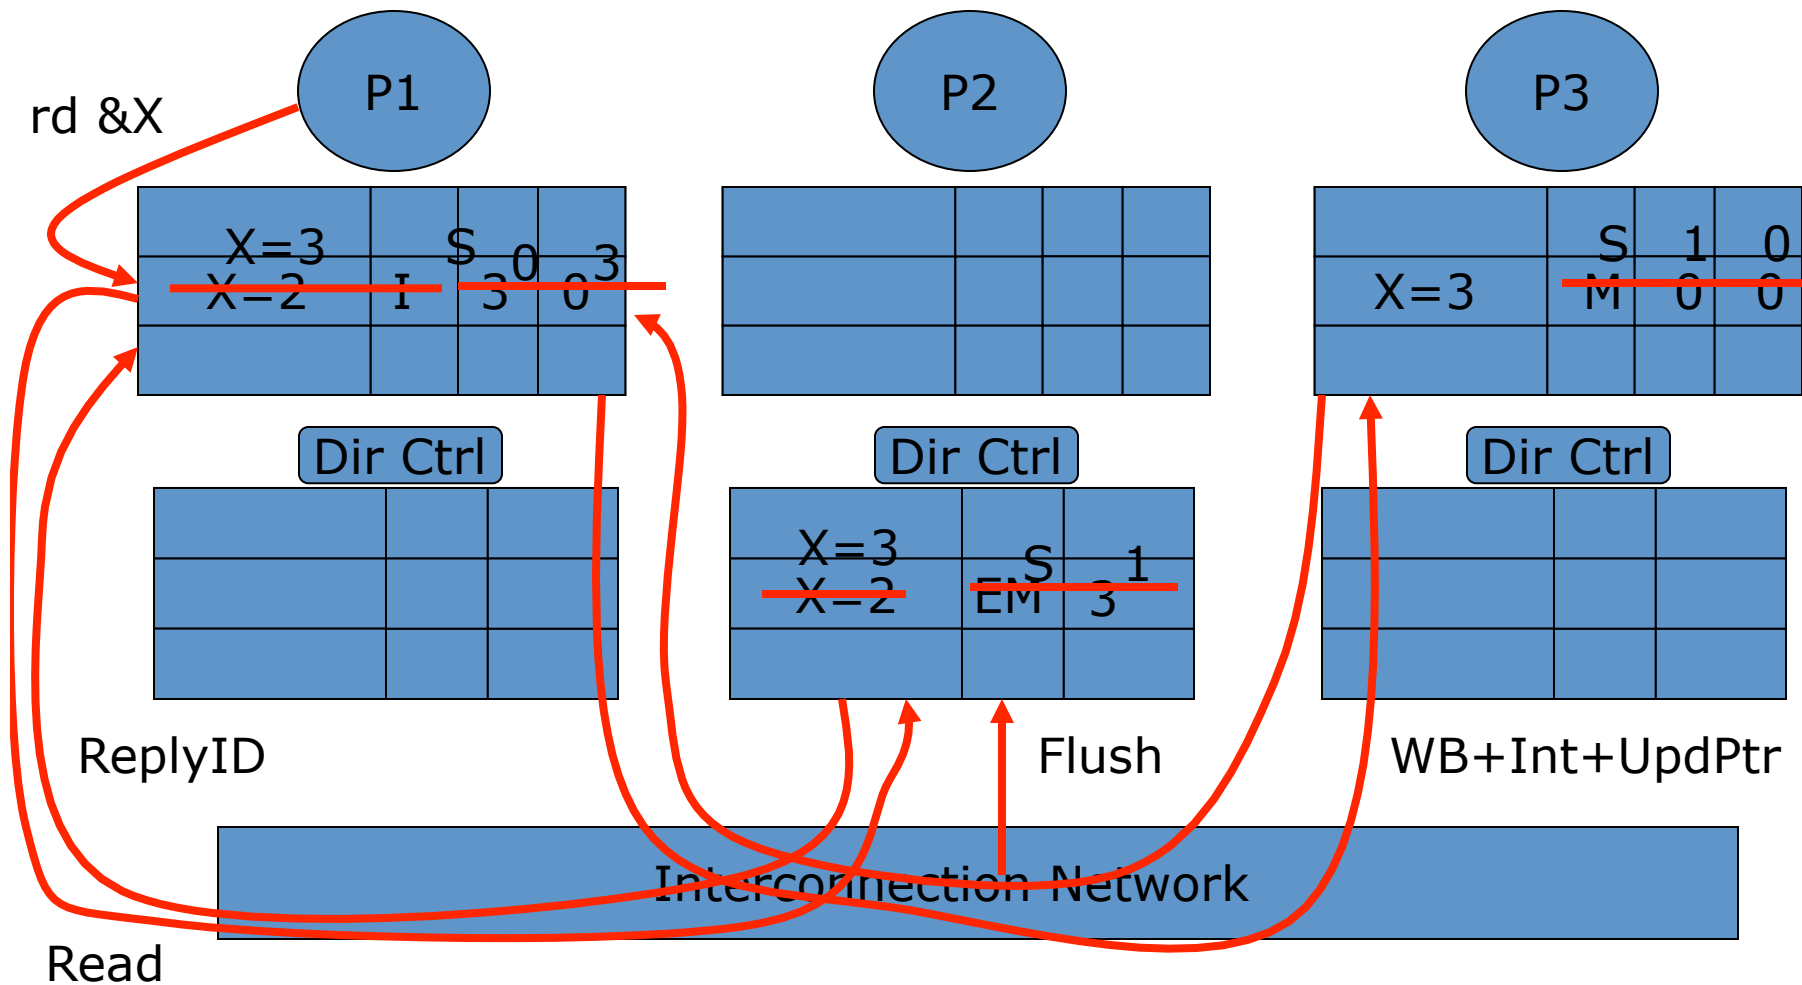
SSCI Protocol Visualization



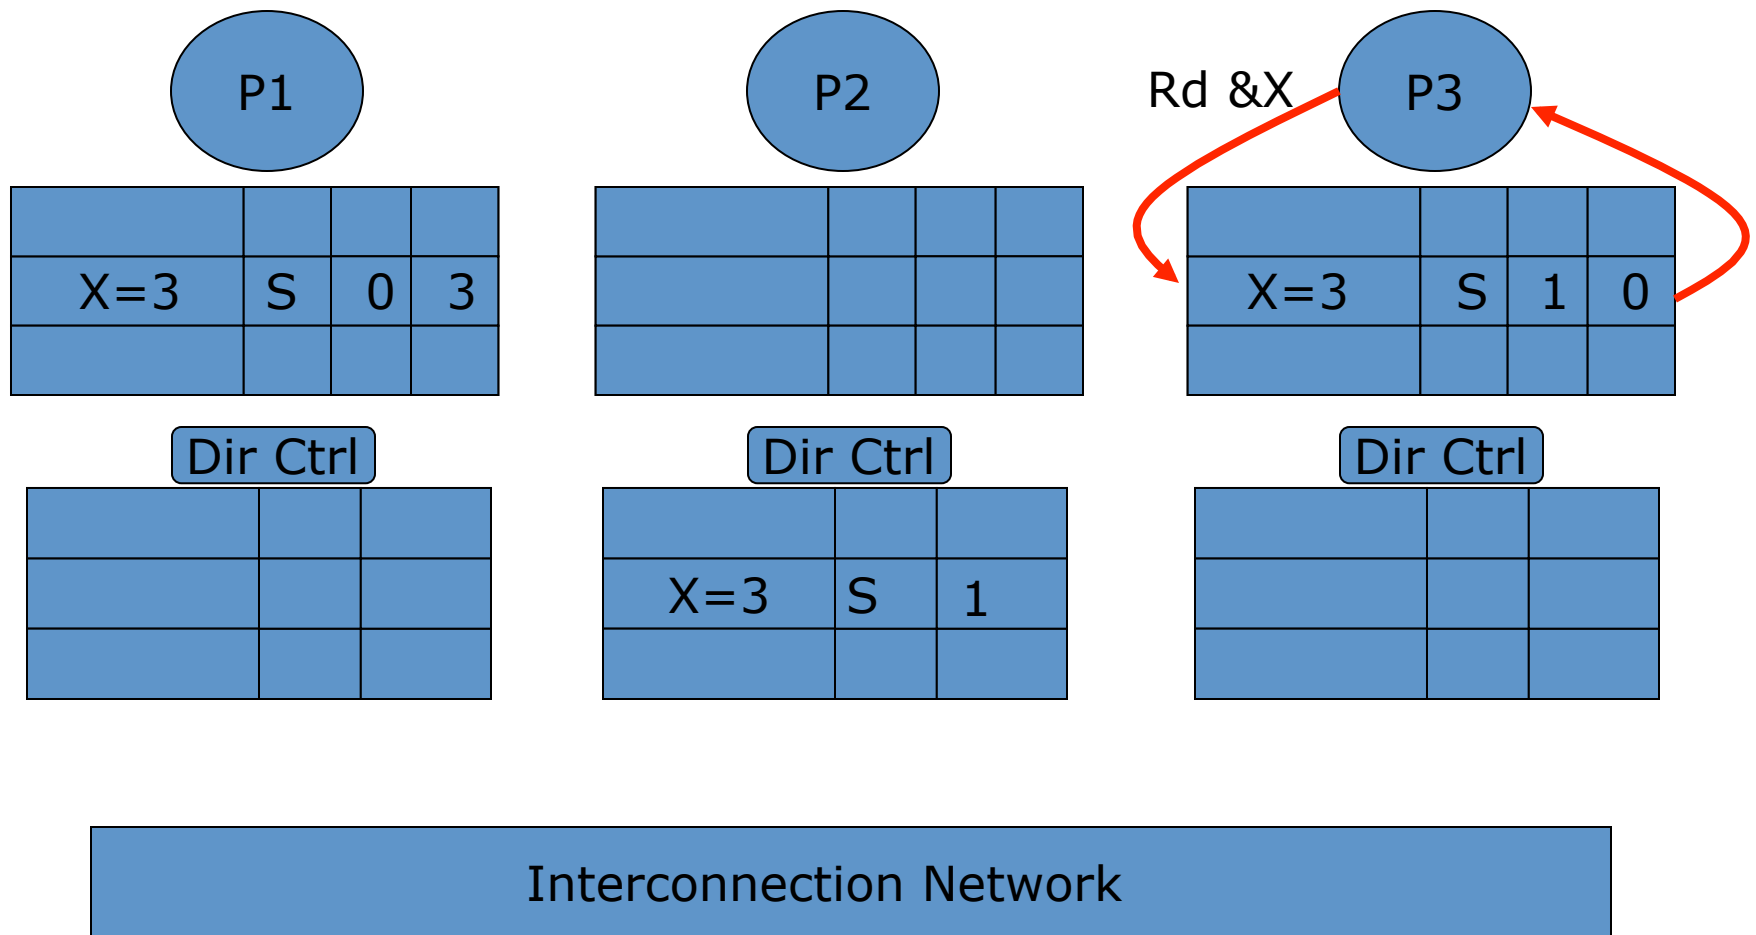
SSCI Protocol Visualization



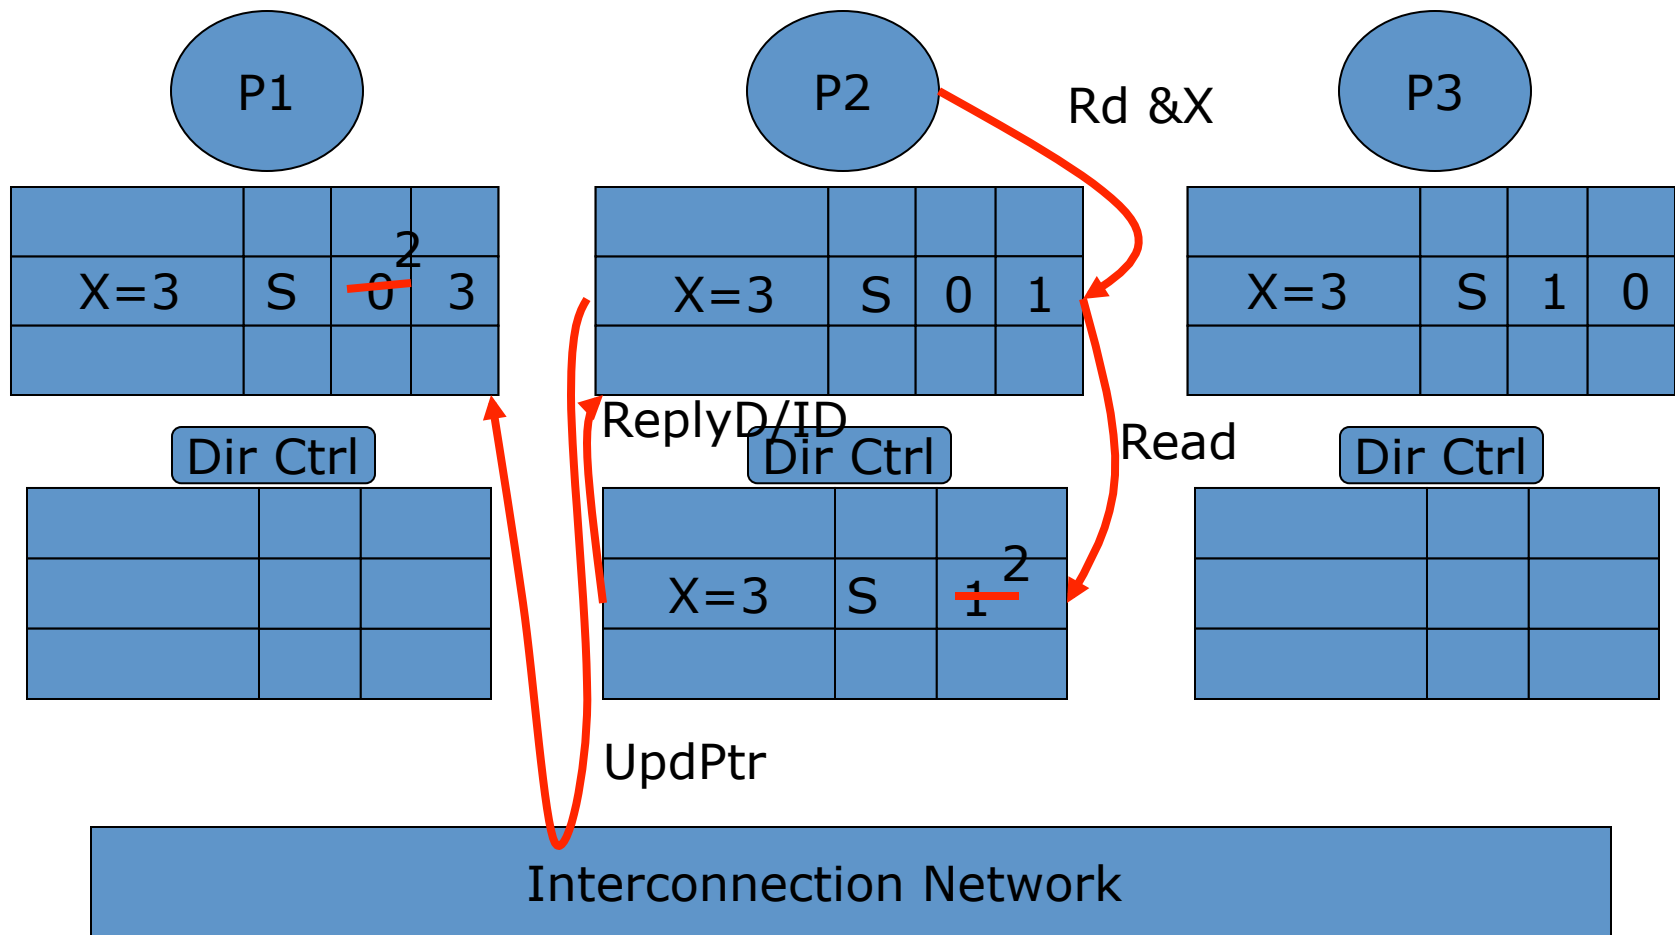
SSCI Protocol Visualization



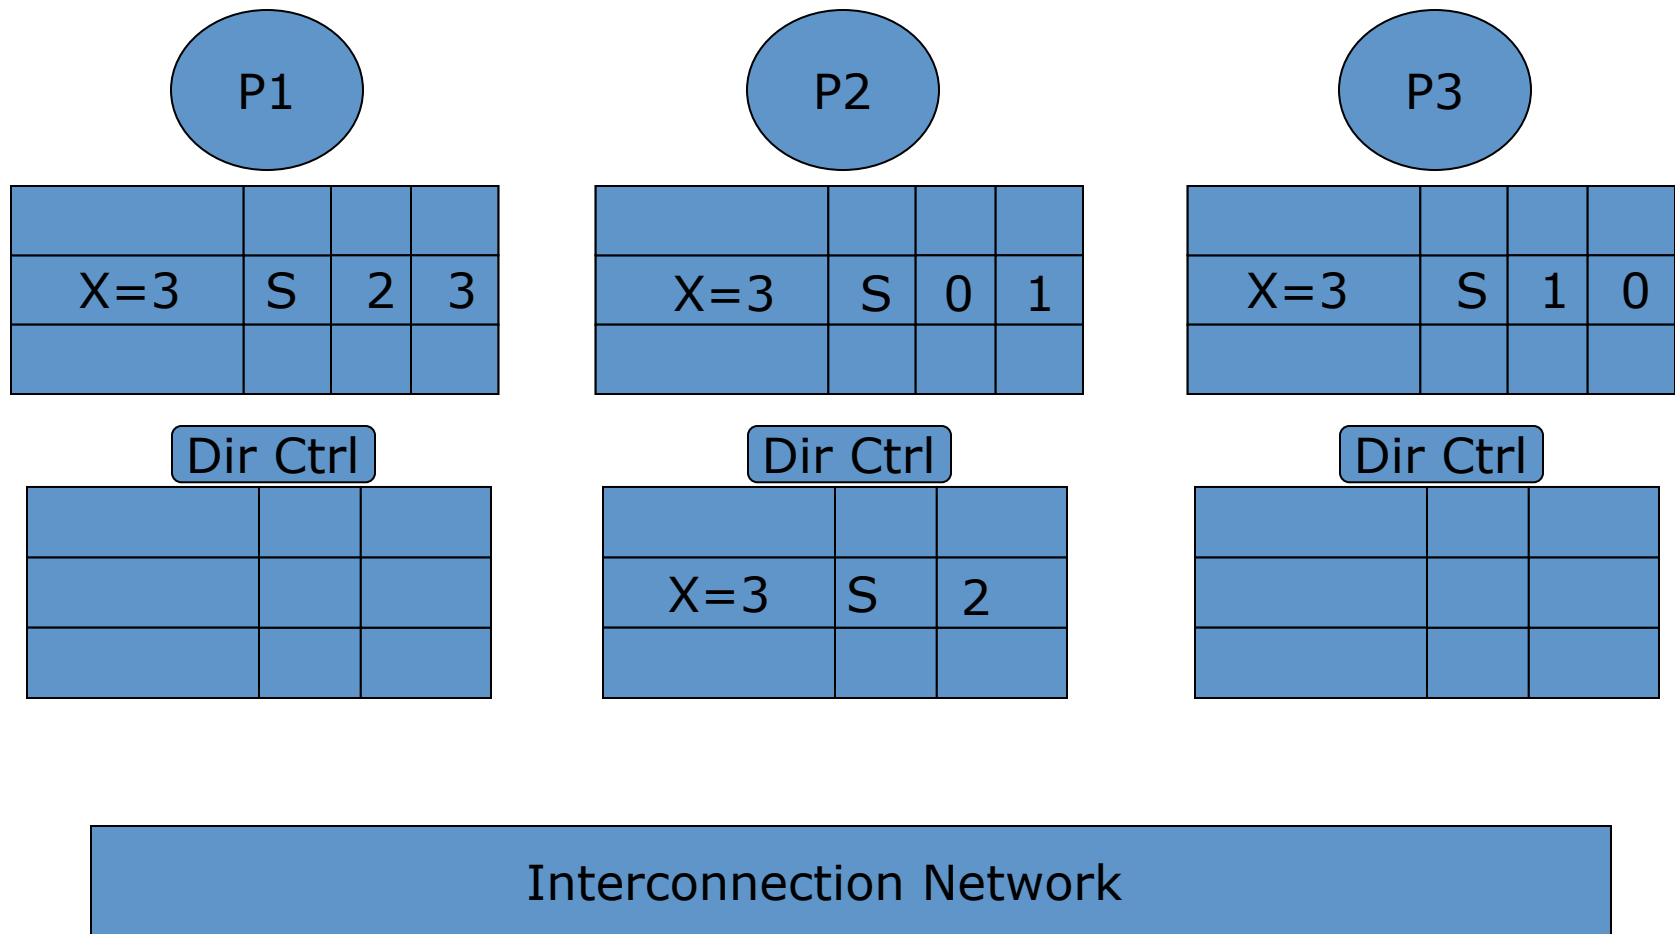
SSCI Protocol Visualization



SSCI Protocol Visualization



SSCI Protocol Visualization



SSCI Example

Proc Action	State P1	State P2	State P3	Dir State @Home	Network Msg	Hops
R1	E,0,0	-	-	EM, 1	Read (P1-> H), ReplyD (H->P1)	2
W1	M,0,0	-	-	EM, 1	-	0
R3	S,3,0	-	S,0,1	S, 3	Read (P3->H), ReplyID (H->P3), WB+Int+UpdPtr (P3->P1), Flush (P1->H, P3)	4
W3	I,3,0	-	M,0,0	EM, 3	Upgr (P3->H) // Inv (P3->P1) InvAck(P1->P3)	2
R1	S,0,3	-	S,1,0	S, 1	Read (P1->H), ReplyID (H->P1), WB+Int+UpdPtr (P1->P3), Flush (P3->H, P1)	4
R3	S,0,3	-	S,1,0	S, 1	-	0
R2	S,2,3	S,0,1	S,1,0	S, 2	Read (P2->H), ReplyD/ID (H->P2), UpdPtr(P2->P1)	3

Observations

- Cache-based vs. Memory-based directory
 - Disadvantages
 - Read to “Shared” data has higher latency and traffic, due to pointer updates
 - Read to “Modified” data needs 4 hops
 - Can be reduced to 3 if home forwards the read request to owner
 - Write invalidation latency is $O(\text{sharers})$
 - Replacement incurs pointer update
 - Advantages
 - Small storage overhead
 - Linked list can be used to support fairness
 - Distributed invalidation
- Protocol race conditions
 - Quite a few more
 - Pointer updates and state transition have to be atomic

Scaling Properties (Cache-based)

- Traffic on write: $O(\text{sharers})$
- Latency on write: $O(\text{sharers})$
 - don't know identity of next sharer until reach current one
 - also assist processing at each node along the way
 - (even reads involve more than one other assist: home and first sharer on list)
- Storage overhead: quite good scaling along both axes
 - Only one head ptr per memory block
 - rest is all proportional to cache size
- Other properties (discussed later):
 - good: mature, IEEE Standard, fairness
 - bad: complex

Comparisons (Fill as Homework)

	Bus-based broadcast	Full bit- vector	Coarse bit-vector	Limited Pointers	Cached dirs (sparse directory)	SCI
Traffic (#msgs on inval)	$O(1)$					$O(\text{sharers})$
Write inval latency	$O(1)$					
Storage overheads	zero	Mp^2 bits				

p = num processors

g = group size (num procs represented by 1 bit)

M = number of memory lines per processor/node

C = number of cache lines per processor

*: $O(p)$ total num cache lines * $O(p)$ bits(procs)/line