

Chapter 12

Interconnection Network

Copyright @ 2005-2008 Yan Solihin

Copyright notice:

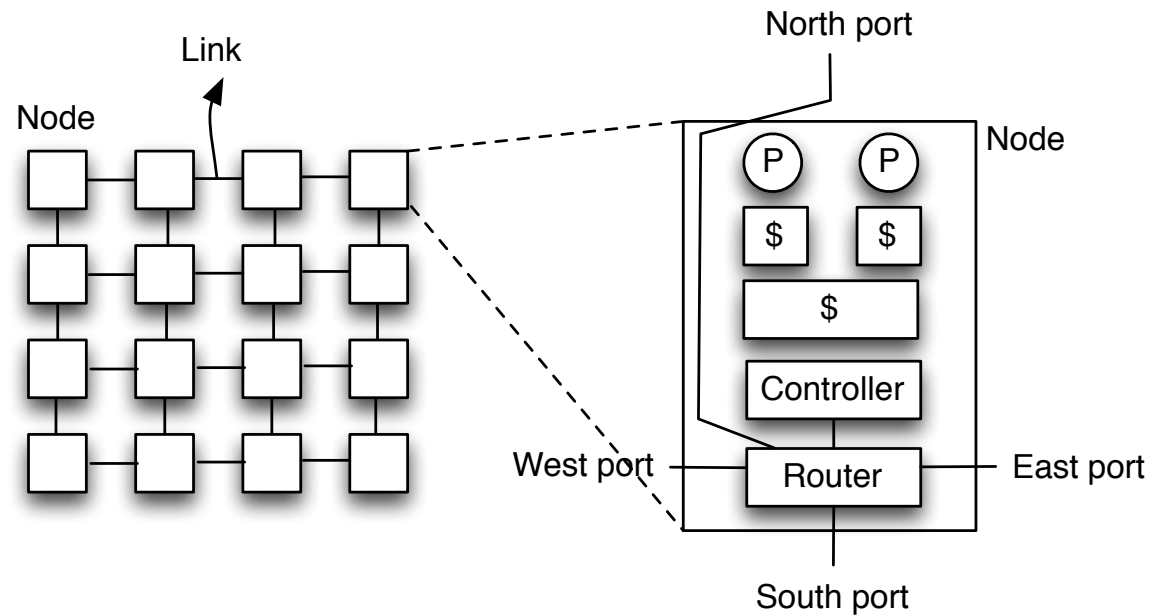
No part of this publication may be reproduced, stored in a retrieval system, or transmitted by any means (electronic, mechanical, photocopying, recording, or otherwise) without the prior written permission of the author.

An exception is granted for academic lectures at universities and colleges, provided that the following text is included in such copy: "Source: Yan Solihin, Fundamentals of Parallel Computer Architecture, 2008".

Introduction


- Interconnection network implements the coherence mechanism that we have discussed so far
 - Sending requests from one node to another
 - Sending data replies from one node to another
- Characteristics of communication in DSM
 - Latency must be low
 - Bandwidth must be high
 - Message characteristics
 - Many small messages with several fixed sizes
 - The largest message is one cache block size (64 or 128 bytes)
- Implications
 - Only two layers are sufficient: link level and node level
 - Communication protocol must be simple (no packet dropping, no elaborate flow control)

Nodes and Links




- A node encapsulates
 - one or more processors + memory hierarchy
 - communication controller
 - Router, which interconnects a node with other nodes

Link and Channel




- Link = set of wires interconnecting a pair of nodes
 - May be unidirectional or bidirectional
 - Unit of transfer: flow control unit (flit)
 - Flit size determined by link width, latency, and amount of buffering
 - Channel = link + sender + receiver
- Flow Control
 - Receiver sends stop or go signal depending on available buffer space
 - Sender reacts by stopping/resuming flit transmission
- A message transmitted over a channel is broken into packets
 - A packet has a header, trailer, and payload
 - Packet broken up into flits



Network Topology

- = shape of the network
- Determines the distance that a message needs to travel (in # links or network hops) from one node to another
- Important characteristics
 - Diameter: largest distance
 - Average distance
 - Degree: how many links are connected to a switch
 - Bisection bandwidth: the minimum number of links cut when the network is partitioned into two equal halves

Popular Network Topologies



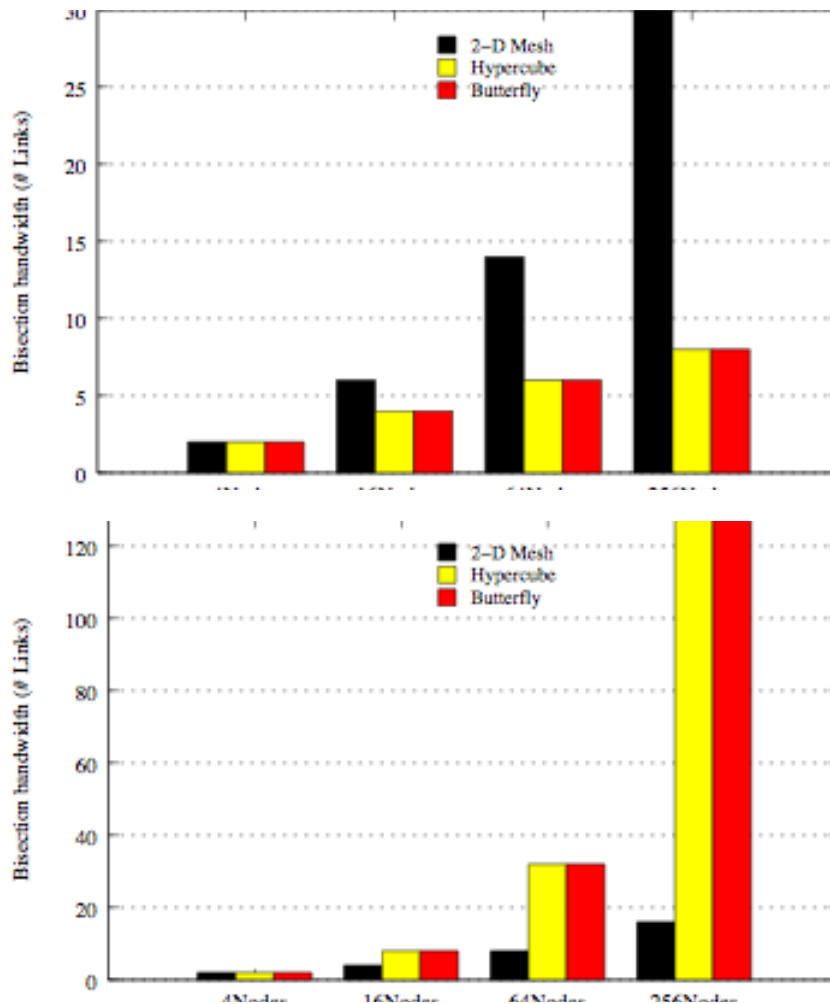
Topology families

- k-ary d-dimensional mesh
 - d = dimensions
 - k = #nodes/dimension
 - Hypercube = binary d-dimensional mesh
- k-ary tree
- butterfly

Comparison of Topologies

Topology	Diameter	Bisection BW	#Links	Degree
Linear array	$p-1$	1	$p-1$	2
Ring	$\frac{p}{2}$	2	p	2
2-D Mesh	$2(\sqrt{p} - 1)$	\sqrt{p}	$2\sqrt{p}(\sqrt{p} - 1)$	4
Hypercube	$\log_2 p$	$\frac{p}{2}$	$\frac{p}{2} \times \log_2 p$	$\log_2 p$
k-ary d Mesh	$d(k - 1)$	k^{d-1}	$dk^{d-1}(k - 1)$	$2d$
k-ary Tree	$2 \times \log_k p$	1	$k(p - 1)$	$k + 1$
k-ary Fat Tree	$2 \times \log_k p$	$\frac{p}{2}$	$k(p - 1)$	$k + 1$
Butterfly	$\log_2 p$	$\frac{p}{2}$	$2p \times \log_2 p$	4

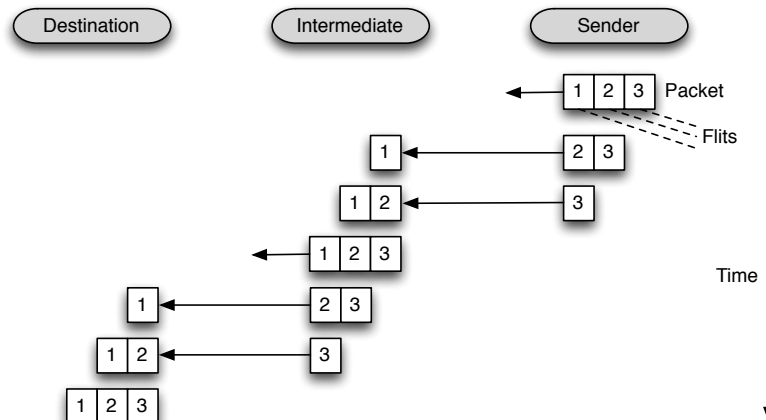
Mesh: Increasing k vs. increasing d




- When more nodes need to be added, is it better to increase the arity (k), or increase the dimension (d)?
 - 2-D mesh vs. hypercube vs. butterfly
- Diameter increases slower in hypercube & butterfly
- Bisection bandwidth increases faster in hypercube and butterfly
- Unfortunately, routers in hypercube are more expensive (higher degrees); while butterfly needs a lot of routers

Routing policy

Store and Forward



Cut-through or wormhole



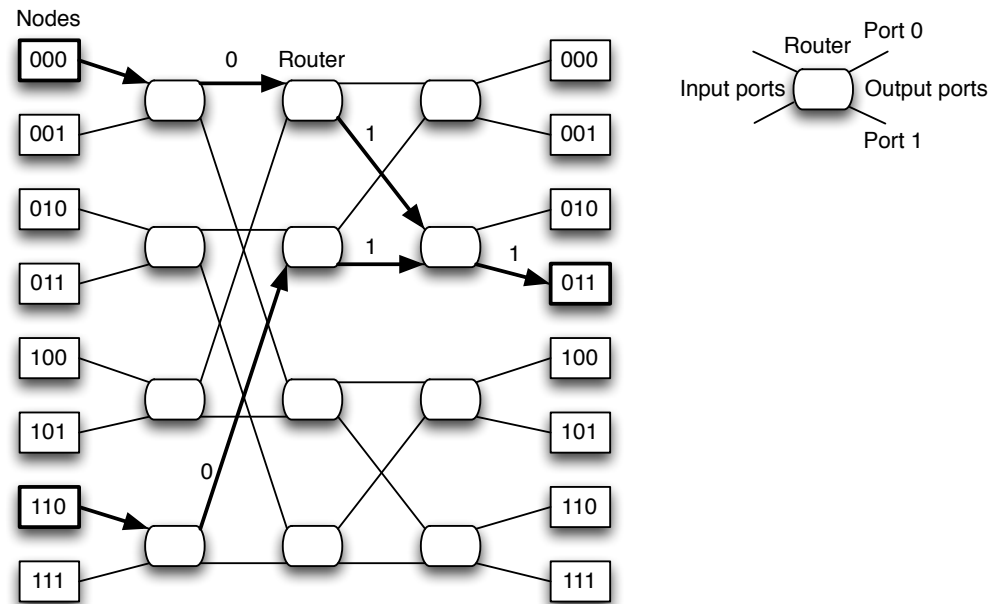
- How to send a packet over the network?
- Store&Forward:
 - Send all flits over a link, then over another, and so on
 - $T = H * (Tr + L/B)$
 - H is distance, Tr is router delay, L is packet length, B is channel bandwidth
- Cut-through:
 - Send flits in pipeline over the links
 - $T = H * Tr + L/B$
- Implication: cut-through routing reduces the importance of topology

Routing Policy


- Path diversity
- Minimal vs. non-minimal
 - Minimal = each packet travels the least number of hops
- Deterministic vs. non-deterministic
 - Deterministic = each sender-dest pair uses a single path for all messages
- Adaptive vs. non-adaptive
 - Adaptive = message can change path as it is being transmitted
- Deadlock-free vs. Deadlock possible
 - Deadlock due to several messages mutually waiting for buffer space

Topology-Dependent Routing

- In some cases, topology imposes routing
- E.g. butterfly: minimal, deterministic, deadlock-free, non-adaptive
 - Lack of path diversity reduces usable bandwidth on adversarial traffic patterns
- Destination ID is directly used to determine output port in each router along the path



Routing in Mesh/Torus



- Dimension-ordered
 - X-Y: Go in X dimension first, and then Y dimension until destination is reached
 - Y-X is possible as well
 - Simple but lacking path diversity


Valiant routing

- Select a random intermediate node
- Can be minimal or non-minimal
- Improved path diversity


Up*/Down* Routing

- Enables routing in any topology
 - Step 1: Start by choosing a root
 - Step 2: Perform breadth-first traversal for next hop, all nodes visited become children, repeat until all nodes are visited. Each link connecting child to parent is an “up” link
 - Step 3: For each level, if two nodes are linked, an “up” link is from a lower-numbered node to a higher-numbered node


Network



Tree construction:
after nodes are visited

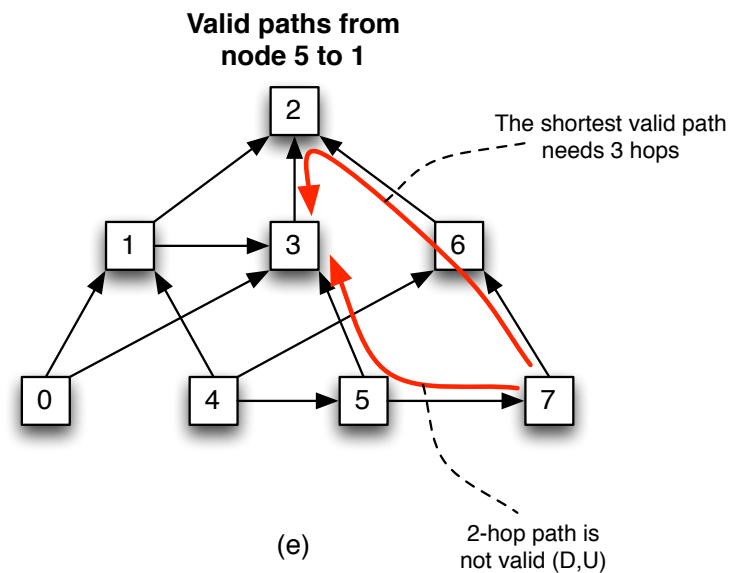



Tree construction:
after all links are visited



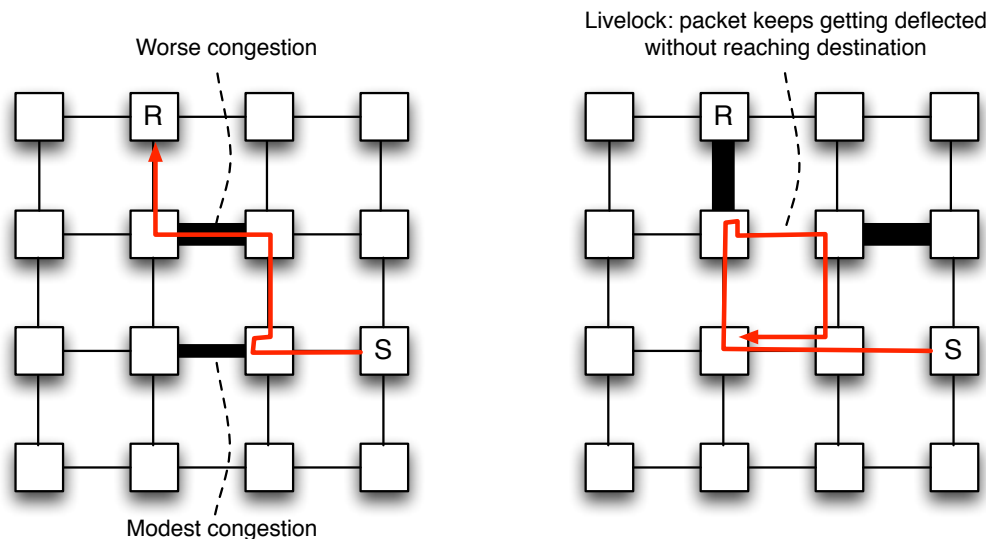
Up*/Down* Routing

- Valid paths may go up first, then down, until destination
 - Down then up is not allowed
- Path diversity is provided, but paths may not be minimal




Adaptive Routing

- Adapts to the state of the network
 - Avoids congestion by deflecting a packet to another port
 - May be minimal or not minimal
 - Risks: (1) May face worse congestion after deflection, (2) possibility of livelock – may want to put a limit on number of deflections




What Deadlock?



- Inability of the network to forward packets
- Due to limited buffer space and cyclic dependence on buffer acquisition
- The figure shows four packets filling up the input buffer and need the next output buffer, but the output buffer is already full
- Handling a deadlock
 - Detect and break: expensive
 - Drop packets: adverse effect in performance and complicates protocol
 - Avoid: restriction on routing
 - Adapt: Escape channel

Turn Restriction in Dimension-Ordered Routing


No turn restrictions:




Dimension-ordered (X-Y) routing:




West First :




North last:



Negative first:




Example



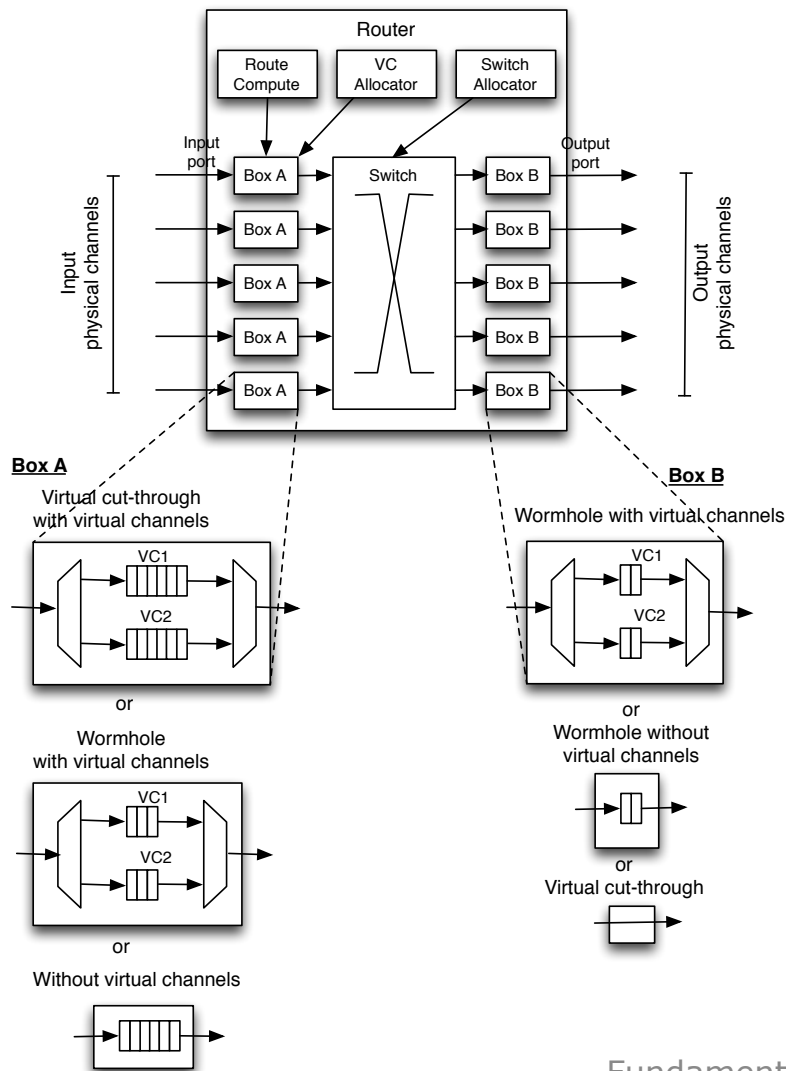
- Allowed in x-y dimension-ordered routing:
 - Path 4
- Allowed in west first routing:
 - Path 2, 3, 4
- Allowed in north-last routing:
 - Path 1, 3, 4

Deadlock

- Avoiding down->up turn in up*/down* routing avoids deadlock
- Use virtual channel to escape from possible deadlock
 - Routing in the escape channel must be deadlock-free



Router Architecture



- Complexity depends on
 - Whether virtual channels are used
 - How many virtual channels share a physical channel
 - Whether wormhole routing or virtual cut-through is used
- Head flit (first flit in the packet) must go through
 - Decode & Compute (DC): decode header, compute output channel, buffer at input channel
 - Virtual Channel Allocation (VA): allocate output virtual channel for the packet
 - Switch Allocation (SA): allocate switch to connect input and output virtual channel
 - Switch Traversal (ST): transmit the flit through the switch to the output virtual channel

